Monitoring of Agents for Dynamic Pricing in different Recommerce Markets

Monitoring von Agenten zur dynamischen Bepreisung
in unterschiedlichen Recommerce-Märkten

Nikkel Mollenhauer

Universitätsbachelorarbeit
zur Erlangung des akademischen Grades

Bachelor of Science
(B. Sc.)

im Studiengang IT-Systems Engineering

eingereicht am 30. Juni 2022 am Fachgebiet
Enterprise Platform and Integration Concepts
der Digital-Engineering-Fakultät der Universität Potsdam

Gutachter	Dr. Michael Perscheid
Betreuer	Dr. Rainer Schlosser
	Johannes Huegle
	Alexander Kastius

Abstract

Sustainable recommerce markets are growing faster than ever. However, businesses now face the challenge of having to price the same item three times: One price for the new item, one for its refurbished version and the price at which items are bought back from customers. Since these prices are heavily influenced by each other, traditional pricing methods become less effective. To solve this dynamic pricing problem, a simulation framework was built which can be used to train artificial vendors to set optimal prices using Reinforcement-Learning algorithms. Before employing these trained agents in real markets, their fitness must be monitored and evaluated, as even the smallest mistake by the agent can lead to high costs for the business. This thesis introduces a number of ways that such dynamic pricing agents can be monitored. We come to the conclusion that using a wide range of tools, from running large-scale simulations to monitoring policy changes following shifting market states, is best when evaluating different aspects of an agent's performance.

Zusammenfassung

Nachhaltige Recommerce-Märkte befinden sich in stetigem Wachstum. Dies stellt Unternehmen jedoch vor die neuartige Herausforderung, dasselbe Produkt mehrfach bepreisen zu müssen: Preise sowohl für die neue und generalüberholte Version sowie ein Ankaufpreis für gebrauchte Ware müssen gesetzt werden. Da diese Preise voneinander abhängig sind, greifen traditionelle Methoden der Preissetzung schlechter. Zur Lösung dieses dynamischen Bepreisungsproblems wurde eine Simulationsplattform gebaut, auf der mithilfe von Reinforcement-Learning-Algorithmen maschinelle Verkäufer für den Einsatz in realen Märkten trainiert werden können. Bevor dies jedoch geschehen kann müssen die trainierten Modelle bezüglich ihrer Eignung überprüft und bewertet werden, da bereits der kleinste Fehler zu hohen Verlusten aufseiten des Unternehmens führen kann. Diese Arbeit führt Tools ein, die für ein Monitoring solcher Modelle verwendet werden können. Wir stellen fest, dass die Nutzung möglichst diversifizierter Methoden, von der Simulation großangelegter Märkte bis zum Monitoring kleinster Verhaltensänderung aufgrund geänderter Marktzustände, die besten Ergebnisse bei der Bewertung einer Bepreisungsmethode liefert.

Contents

1 Introduction

This thesis builds upon the bachelor's project 'Online Marketplace Simulation: A Testbed for Self-Learning Agents' of the Enterprise Platform and Integration Concepts research group at the Hasso-Plattner-Institute. Therefore, the project will be referenced and all examples and experiments will have been conducted using its framework.

1.1 Motivation

Nowadays, shoppers and retailers alike are getting more and more environmentally conscious. A study conducted in 2020 found that over two thirds of shoppers planned on buying sustainable clothing in the future, and over half already did buy sustainable clothing regularly [KPM20]. At the same time another study reveals that retailers favour online channels over offline channels when selling used goods, with 78% preferring the former and only 6% the latter [Ins20]. This demand is leading more and more companies, especially those selling their products through e-commerce channels, to adopt more sustainable strategies and enter the *Circular Economy*, a concept we will introduce in **1.2 The Circular Economy model**. In the context of e-commerce, Circular Economy markets are also referred to as *recommerce* markets, a phrase coined in 2005 [Arc05].

The goal of the bachelor's project this thesis is based on was to simulate such recommerce markets, to allow for risk-free and efficient learning of market characteristics for different pricing agents. Aside from classically Rule-Based pricing methods, the project also focussed on training machine learning models using *Reinforcement-Learning* on the simulated marketplace, see **1.3 Reinforcement-Learning**. The focus of this thesis lies on ways that such dynamic pricing agents can be monitored and evaluated. In **2 Related Work** we will first explore the difficulties faced in dynamic pricing, as well as current and novel approaches to monitoring and evaluating Reinforcement-Learning agents in different scenarios. This will be followed by an overview of the specific features of the *recommerce* market that we implemented in **3 Simulating the marketplace**. In **4 Approaches to monitoring agents**, we will give a detailed explanation of the different tools we built and used to monitor and evaluate our different agents. These tools will be put into context in **5 The *recommerce* workflow**, where the different parts of the framework will be put to use. Finally, in **6 Analysing Graphs**, we will train a Reinforcement-Learning agent using our simulation framework and take a look at how the different monitoring tools can help in evaluating its performance.
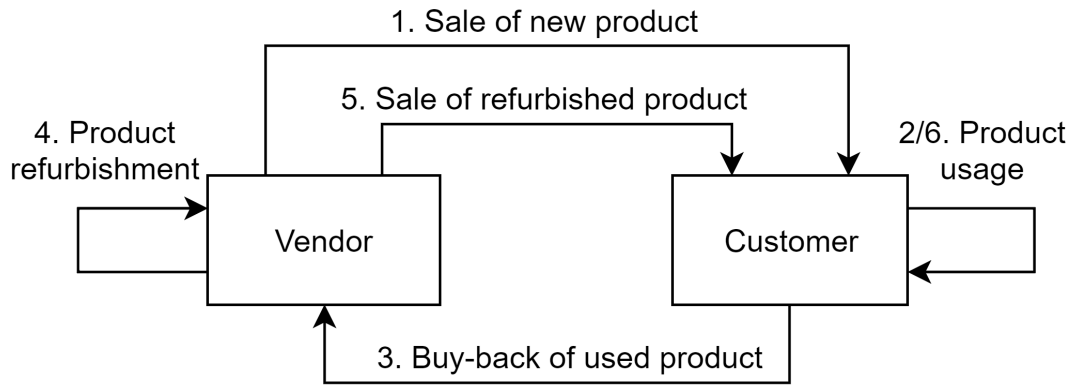
Figure 1.1: The product lifecycle in a Circular Economy. In a Linear Economy, the product lifecycle ends with step No. 2.

1.2 The Circular Economy model

The main goal of the aforementioned bachelor's project was to develop an online marketplace that simulates a realistic Circular Economy market environment. A market is most commonly referred to as being a *Circular Economy* if it includes the three activities of reduce, reuse and recycle [KRH17]. This means that while in a classical Linear Economy market each product is being sold once at its *new price* and after use being thrown away, in a Circular Economy, a focus is put on recycling and thereby waste reduction. In our project, we first started by modelling the simpler Linear Economy, upon which we then built the more complex Circular Economy markets. This was done by adding two additional price channels, *used price* and *rebuy price*, to the pre-existing *new price* of a product. Refer to Figure 1.1 for an overview of the product lifecycle in a Circular economy.

The *rebuy price* is defined as the price a vendor is willing to pay a customer to buy back a used product, while the *used price* is defined as the price the vendor sets for products they previously bought back and now want to sell alongside new products (whose price is defined by the *new price*). In **3.1 Market scenarios**, we will explain the different market scenarios we modelled in more detail, together with how we transferred different aspects of these scenarios from the real market to our simulation framework. Please also refer to Figure 3.1 for an overview of the different market components and how they interact in our framework.

From now on, when mentioning a general *market* or *marketplace*, we are referencing a Circular Economy marketplace with rebuy prices.

1.3 Reinforcement-Learning

After the initial market was modelled the goal was to train agents using different Reinforcement-Learning algorithms to dynamically set prices on this marketplace, both in monopolistic scenarios and in competition with Rule-Based vendors which set prices following a strict set of pre-defined rules. These rules can range from simply undercutting the lowest competitor's price to more advanced techniques such as smart inventory management and reliance on previous sales data. An overview of the different types of vendors, both Rule-Based and using Reinforcement-Learning can be found in **3.3 Vendors**. Furthermore, functionality was added

Figure 1.2: The standard Reinforcement-Learning model in the context of our market simulation.

that allows for different Reinforcement-Learning algorithms to be trained against each other on the same marketplace, as well as functionality for so-called *self-play*, where an agent plays against itself, or more precisely, against its own policy.

Reinforcement-Learning agents are trained through a process of trial-and-error. They interact with the market through an observable state and an action which influences the following state. Figure 1.2 illustrates the Reinforcement-Learning model in the context of our market. The goal of the agent is to maximize its *reinforcement signal*, which in the case of our simulation framework is the profit the agent made during the last episode, since we want to train agents to maximize profits on real markets. An episode consists of a fixed, but pre-configurable number of time steps, where in each step each vendor (agent) sets their prices and customers make purchasing decisions. In this sense, a simulated episode could be imagined to be a day in the real world, with vendors updating their prices multiple times per day, whenever a new time step begins. By observing which prices lead to which profits (reinforcement signal), the agents get more effective in their pricing decisions over the course of training, which spans thousands of episodes in most cases.

2 Related Work

This section will outline approaches to modelling and simulating recommerce market-places, as well as give an introduction into the concepts of Reinforcement-Learning, the technology the framework is built for. Additionally, novel and unconventional approaches to monitoring and evaluation of these Reinforcement-Learning agents will be presented.

2.1 Dynamic Pricing simulations

The topic of dynamic pricing techniques is well explored, with the earliest mathematical models having been developed over a century ago [den15]. With the emergence of e-commerce and the ability to cheaply and quickly change prices, as well as the freedom of information, including being able to know competitor's prices in real time, the topic has gained importance even further. However, research concerning dynamic pricing in e-commerce, especially in highly competitive markets, has not grown at the same pace [GB22]. On the other hand, the role that autonomous agents will play in this new market environment, both in the form of vendors and customers, has for a long time been a topic of discussion and speculation alike [KHG00]. Dynamic pricing methods come in many different shapes and forms, from simple greedy algorithms over customer-choice methodology to machine-learning models [Gee+19]. Due to this high number of options to choose from, and to enable researchers to evaluate their algorithms' performance, a unified platform for comparison is needed. The real market is not really an option for this, as it is near impossible to create equal opportunities for each pricing method to be able to reliably compare them to each other. This is one of the major reasons for simulating the real market - development, testing and comparison of new pricing methods.

2.2 Visualization - State-of-the-art

The process of training Reinforcement-Learning agents for any kind of task always comes with the requirement for visualizing the data collected during training. This allows for an analysis of an algorithm's performance and gives insights into its strengths and weaknesses.

For the past years, going back as far as 2018, one of the most used frameworks overall and the most used for machine learning was TensorFlow [Ove22]. Aside from its API for model building, TensorFlow also provides a front-end visualization toolkit called TensorBoard, which can be used independent of other TensorFlow tools. TensorBoard provides an API for tracking and visualizing important metrics such as loss and accuracy, and allows developers to easily integrate their own metrics as well. During an experiment, data can be visualized live during the training process, allowing developers to quickly gain insights into the performance of the

algorithms. In our recommerce market simulation framework, we use the Tensor-Board in conjunction with our own tools for data visualization, see **4 Approaches to monitoring agents**.

2.3 Visualization - Novel approaches

2.3.1 Model visualization

Aside from visualizing results of a training run, developers may also want to visualize the model itself. Tools such as the *Graph Visualizer* [Won+18] are a step in the right direction, but the study found that while developers are satisfied with the visualization tool itself, they would prefer to be able to edit and influence the graph model directly. The *SimTF* tool developed by Leung et al. [CLH18] attempts to solve this problem. The same as the *Graph Visualizer*, the *SimTF* tool is based on TensorFlow. The authors describe it as a *library for neural network model building*. *SimTF* acts as a middleman between the visually constructed neural network graph model and the TensorFlow API, allowing developers to modify the visualized model to influence the training network.

In our simulation framework, the neural network model itself is quite static, users can only choose from a pre-set selection of network sizes, and all changes to hyperparameters must be done through configuration files ahead of a training run.

2.3.2 Evaluation

On the other side of the visualization tools we have those which provide insights into currently running or completed training runs, such as the previously mentioned *TensorBoard*. Besides simply visualizing different collected data, the goal of most of these tools is to allow the developers to get a sense of the trained algorithm's performance, and to evaluate and compare it against previously trained agents using the same algorithm or other algorithm's completely.

When developing new algorithms, the most common barrier to proper comparison with current state-of-the-art algorithms is a lack in reproducibility of results. Benchmark environments such as those provided by OpenAI Gym [Bro+16] lower this barrier by providing unified interfaces against which many algorithms have already been tested. However, the effects of extrinsic factors (such as hyperparameters) and intrinsic factors (such as random seeds for environment initialization) make proper, reliable reproducibility a challenge, something for which current evaluation practices do not account for [Hen+17]. Jordan et al. [Jor+20] propose a new, *comprehensive evaluation methodology for reinforcement learning algorithms that produces reliable measurements of performance*. The authors describe the goal of their new evaluation technique as not trying to find methods (algorithms) that maximize performance with optimal hyperparameters, but rather to find those that do not require extensive hyperparameter tuning and can therefore easily be applied to new problems. This leads to a preference for algorithms that are not aimed at being the best at problems which are already solved (which applies to the aforementioned benchmark environments), but instead those which are most likely to succeed in new, unexplored environments.

3   Simulating the marketplace

*This chapter will introduce the various components that make up a recommerce market, and how they were implemented in our simulation framework. We will take a brief look at different market scenarios as well as how our customers make decisions. The focus of this chapter will however lie on the vendors, the Rule-Based as well as the ones trained using Reinforcement-Learning algorithms, comparing their features and how they fare against each other. These comparisons will be done using our various monitoring tools, which will be explained in the following chapter, **4 Approaches to monitoring agents** and put to use in **6 Analysing Graphs**. For further information on the overall framework structure refer to [Dre22], for detailed insights into our market processes, see [Bes22].*

3.1 Market scenarios

In our framework, we implemented a number of 'market blueprints' for classic Linear Economy markets as well as for Circular Economy markets both with and without rebuy prices enabled. There are marketplaces available for each combination of the following two features:

1. Marketplace type: Linear Economy, Circular Economy, Circular Economy with rebuy prices

2. Market environment: Monopoly, Duopoly, Oligopoly

The marketplace type defines the number of prices the vendor has to set. In a Linear Economy, vendors only set prices for new items, in a Circular Economy prices for refurbished items need to be set as well. In order to simulate a recommerce market, the user can choose a Circular Economy with rebuy prices. The market environment defines the number of competing vendors in the simulation: Monopolistic and competitive markets are available, where the Duopoly is simply a particular version of an Oligopoly. Depending on the chosen market environment, one, two, or any number of vendors can be chosen to be used in the experiment. Figure 3.1 shows a reduced overview of the classes in our framework which concern themselves with the market simulation, and how they interact with each other during the simulation.

In the most common use case of our framework, the training of Reinforcement-Learning agents, only the agent (or multiple agents in the case of training multiple Reinforcement-Learning agents against each other) that is to be trained needs to be configured by the user, as each market environment is equipped with a pre-defined set of competitors that will play against the agent defined by the user. To allow for more control over the simulation, users are however also able to change these competitors to use any vendors they want - as long as they are a valid fit for the marketplace type and the number of competitors chosen matches the market environment.
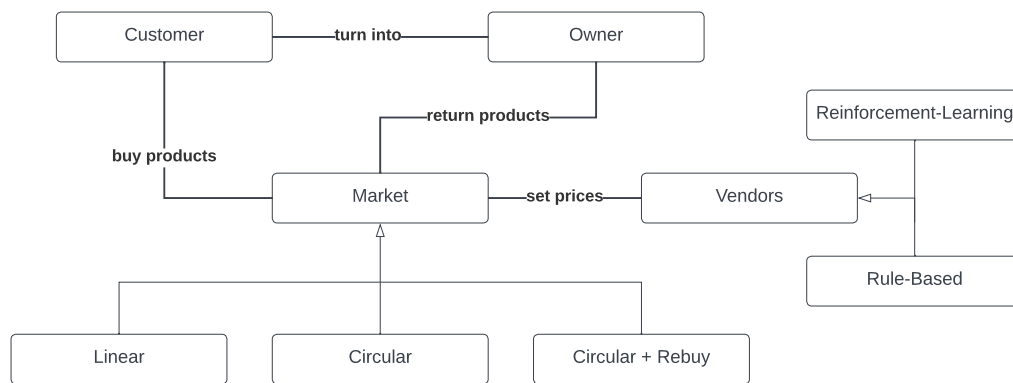
Figure 3.1: Class diagram outlining interactions between classes concerning the market simulation.

3.2 Customers

Customers are at the centre of every type of marketplace, which makes them an integral part of our market simulation. However, since each customer in the real market is an individual with different backgrounds and makes purchase decisions based on personal preferences, modelling a realistic depiction of real-world customers proves to be very difficult and time-consuming. For this reason we decided to keep our initial implementation of the customers as simple as possible, taking into account future extension and scalability concerns.

Most customers' behaviour can be classified into one of a (non-exhaustive) number of categories, such as *Perfectionistic* or *Impulsive* consumers, proposed in [EIT13] (also see Table A.1). As we are focussing on dynamic pricing and our vendors can only change/influence the prices of their products, we decided on primarily building customers that value price over any other features a product may have, thereby incentivizing our vendors to make the most of their pricing power. This behaviour coincides with the shopping style of the *Price Conscious* or *Value for money* consumer.

To make Linear markets a little more complex and add another dimension than just the *new price* of a product for vendors to consider, random quality values are assigned to each vendor's products. Customers in this economy model were modelled to take this quality value into account when making their purchasing decisions, further reinforcing the shopping style mentioned above. As Circular Economy markets are inherently more complex than their linear counterparts (through the addition of two new price channels, which influence each other through their effect on customers), it was decided to remove the additional layer of quality values from these markets.

Within each simulated step, after the various vendors have set their prices, purchase decisions are made by the customers. To save on computational time, probability distributions are used to determine what part of the total number of customers will decide to take which action. See below for the way these distributions are calculated for an exemplary Circular Economy marketplace, for further information also see [Bes22].

Number of pages: I could remove this next paragraph, since we focus on Circular Economies

▶ **Definition 3.1.** Let P_{in} be the price of the new and P_{ir} the price of the refurbished product of vendor i. Using these prices, for each vendor, we define R_{in} as the *preference ratio* regarding P_{in}:

$$R_{in} := \frac{10}{P_{in}} - e^{P_{in}-8}$$

and similarly, R_{ir} as the *preference ratio* regarding P_{ir}:

$$R_{ir} := \frac{5.5}{P_{ir}} - e^{P_{ir}-5}$$

◀

In order to be usable as a probability distribution, the different ratios need to add up to 1. We use *softmax* to normalize our values, after which each *preference ratio* will be within the interval $(0, 1)$, with all ratios adding up to 1 as needed.

Following this, the market uses a *multinomial* distribution to draw n samples, with n being the number of customers making a pruchasing decision in this step of the simulation.

As mentioned above, the current decisionmaking process of customers in our simulation is still quite basic. **7.1 Modelling a realistic *recommerce* marketplace** introduces a number of parameters and circumstances that can be used to make customer behaviour more realistic. Through the modular approach when building the framework, any customer behaviour implemented in the future can easily be added to the pool of available options, as long as it manifests in the form of a probability distribution, at which point the number of customers can be split between different distributions when drawing from the *multinomial* distribution (or any other distribution if so desired).

@Rainer @Alex: Should I add a graph for the preference ratios in the Appendix? AFAIK, Nick has one in his thesis, so I could also just leave it, as I already cited him.

3.2.1 Owners

Once a customer has decided to buy a product from any of the available vendors, they turn into an *owner*. In the next step of the simulation, they are offered the option of selling their now used product back to one of the vendors. If they decide to do so, the vendor pays them the advertised *rebuy price* and adds the used product to their inventory, from where it can then be sold as a refurbished product in the next step.

In our simulation, all owners are memoryless, meaning that they do not remember the original price they paid for the product. Additionally, each vendor in the market is obligated to buy back any product, independent of the original vendor. In each step, every owner has the option to either keep the product for another step, discard it (meaning it is removed from circulation and not sold back to a vendor), or sell it to any one of the vendors in the market. Similarly to the way we compute customer purchasing decisions, the decisions owners take are also represented through probability distributions. When deciding if and to which vendor a product should be sold, the following formula is used:

▶ **Definition 3.2.** We keep the variables defined in Definition 3.1. Additionally, let P_{ib} be the price at which vendor i is willing to buy back an owner's product. For each vendor, we define P_{im} as the purchase option with the lowest price:

$$P_{im} := \min(P_{in}, P_{ir})$$

We define the likelihood R_{ib} that the owner will return their product to vendor i as follows:

$$R_{ib} := 2 \cdot e^{(P_{ib} - P_{im})/P_{im}}$$

Additionally, the owner's *discard preference R_d* is updated for each vendor as follows:

$$R_d := \min(R_d, \frac{2}{P_{ib} + 1})$$

Meaning the owner is more likely to discard their product if rebuy prices are low across the board. ◀

3.3 Vendors

Vendors are the main focus of our market simulation. While our framework will not be able to reproduce all types of pricing models used in the real market, we strive to model as many different models as possible (and feasible in the scope of the project). Due to the modular nature of the framework, it is possible for users to easily create and add their own pricing strategies in the form of a vendor that can play on a market. This allows users to not only add other Reinforcement-Learning algorithms, but also to define new Rule-Based strategies. For this reason, we also do not restrict the usage of our monitoring tools to just Reinforcement-Learning agents, but allow users to monitor any combination of vendors.

@Rainer @Alex: Do I need to further motivate *why* we allow this?

We will use four types of dynamic pricing models, as defined in 'Dynamic pricing models for electronic business' [Nar+05], to describe the different models we implemented in our framework. These categories are not mutually exclusive, and many agents belong to more than one category.

3.3.1 Inventory-based models

These are pricing models which are based on inventory levels and similar parameters, such as the number of items in circulation (items which are currently in use by customers). In our framework, almost all Rule-Based agents consider their inventory levels when deciding which prices to set. The only exception to this rule are the simplest of our agents, the *FixedPriceAgents*, which will always set the same prices, no matter the current market state and competitor actions. The prices these agents will set are pre-determined through the user's configuration of the experiment, and will not change over the course of the market simulation.

Inventory-based models are comparatively easy to implement, as they only depend on data immediately available to the vendor. This has the advantage that Rule-Based agents which fall into this category are relatively simple to create and modify. Examples of *Inventory-based agents* in our framework can be found in the *RuleBasedCERebuyAgent*, one of the first Rule-Based agents we created. This agent does not take pricing decisions of its competitors in the market into

account, but simply acts according to its own storage costs, always trying to keep a balance between having enough (refurbished) products to sell back to customers and keeping storage costs low. It does this by checking into which of four possible ranges the current number of stored products falls. The more products the agent already has in storage, the lower it sets the rebuy price (to 'prevent' customers from returning more products), and the lower it sets the price for refurbished products, to empty the storage more quickly and prevent high storage costs. While its performance is not necessarily bad, it is still one of the weakest competitors currently available in the framework. For the full implementation logic of the agent's policy, see Figure A.1. For a comparison of the *RuleBasedCERebuyAgent* with a more sophisticated *Data-driven model* see Figure 6.8.

3.3.2 Data-driven models

Data-driven models take dynamic pricing decision-making one level further. They utilize their knowledge of the market, such as customer preferences, past sales data or competitor prices, to derive optimal pricing decisions. Aside from the aforementioned *FixedPriceAgents* and the basic *RuleBasedCERebuyAgent*, all of our other Rule-Based agents fall into this category. One of the most prominent examples of a *Data-driven model* is the *RuleBasedCERebuyAgentCompetitive*, an agent whose basic goal is to always try and undercut competitor prices, while also keeping track of the current amount of items in storage. In its policy, this agent first always undercuts the new price of all other competitors. Then, similar to the *RuleBasedCERebuyAgent*, the agent looks at the current amount of stored products and depending on the range, sets lower prices for refurbished items and higher rebuy prices, while always trying to give customers a better deal than any of the competitors. Again, the full implementation logic of this vendor's policy is available in the Appendix, under Figure A.3.

Notably, all of our *Data-driven models* are also *Inventory-based* to a certain extent, as handling storage plays a big part in a Circular Economy market setting, where used products need to be bought back from customers and subsequently undergo refurbishment while in inventory of the company. *Data-driven models* have proven to be the most competent Rule-Based agents in our recommerce market scenario, in particular the above described *RuleBasedCERebuyAgentCompetitive*, which is able to outperform most other Rule-Based agents, both *Inventory-based* and *Data-driven models* (please refer to e.g. Figure A.7).

3.3.3 Game theory models

Game theory concerns itself with the study of models for conflict and cooperation between rational, intelligent entities [Mye97]. It is therefore often applicable in situations where competing individuals, acting rationally and selfishly, interact with each other. In our framework, most competitors in the market are influenced in their decision-making processes by the actions of other participants of the scenario. This is especially true for the Reinforcement-Learning agents, which base their policy on the received market states, which include their competitor's actions. While none of our Rule-Based agents have been specifically designed to act according to game theoretic strategies, due to the fact that almost all of them consider their competitors prices in their pricing decision, and due to the

> Number of pages: If any section needs to be removed here (e.g. for page length), this should be the one

nature of Reinforcement-Learning trying to maximize their own profits without regard to their competitors performance, behaviour according to Game theory can sometimes be observed. This can be observed especially well in the beginning of training sessions, where Rule-Based competitors often struggle to achieve high profits while the Reinforcement-Learning agent is making great losses, but are then able to increase their profits as the trained agent starts to make better decisions as well. Many examples of such behaviour can be found in the diagrams shown in **6 Analysing Graphs**, such as in Figure 6.1.

During training, Reinforcement-Learning agents observe the market state, which includes prices and sales data not only of themselves but also the other vendors in the market. Using this data, the algorithms try to predict how their prices will influence customer behaviour as well as the competing vendors. In some cases, depending on the concrete behaviour of the competitors, vendors may cooperate in driving prices higher together, in other cases the agents may act seemingly irrationally, lowering prices in order to force their opponents to lower theirs as well. An example of such behaviour can be seen in the development of prices for new items set by the two competing vendors in Figure 6.3 (d).

3.3.4 Machine learning models

All of our Reinforcement-Learning agents fall into this category. As they are not the focus of this thesis, we will not go into detailed explanations of the various models used. For a comparison of the performance of different Reinforcement-Learning algorithms in the context of our simulation framework, please refer to [Gro22] The following will give a short overview over the different algorithms used in our framework.

There are two 'origins' of the algorithms in our framework. In the earlier phases, Q-Learning and Actor-Critic algorithms were custom implemented by us. Later on, we used the Stable-Baselines library to incorporate a greater number of pre-implemented algorithms into our framework. While these are not as easily configurable, they can quickly be used without much additional work.

- *Q-Learning*: *Q-Learning* agents were the first Reinforcement-Learning agents we introduced in our framework, as its algorithm is one of the easier ones to implement [KLM96]. The *Q-Learning* algorithm used in our framework is implemented using the *pytorch* framework (see [Pas+19]). However, the drawback of using *Q-Learning* in our framework is that it can only be applied to discrete action and state spaces. This means that when using *Q-Learning*, only 'whole' prices can be set by our vendors, and any decimal places must be omitted. This of course limits the framework in its realism, as the fine-tuning of prices using decimal places can be critical in influencing customer decisions. In the initial exploration-phase of our simulation framework this was not a problem, as relatively small action spaces were used, adapted to this limitation. But by now, our simulation framework also supports continuous action and state spaces, to allow more complex algorithms such as those introduced in the sections below to function. While approaches for *Q-Learning* algorithms that can work with continuous action and state spaces have been presented in the past ([GWZ99], [MPD02]), we have chosen not to implement such an algorithm in our framework, but rather explore approaches other than *Q-Learning*, such as *Actor-Critic* algorithms, introduced below.

- *Actor-Critic*: *Actor-Critic* algorithms are more complex than Q-Learning algorithms, and have therefore been implemented later in the process. They are structured different then Q-Learning algorithms in the way that they are 'split' into two parts: The *actor* is responsible for selecting actions, while the *critic* is responsible for critiquing the actions taken by the *actor*, thereby improving its performance [SB18]. Similar to the *Q-Learning* agents, the *Actor-Critic* algorithms have also been implemented by us. In total, one discrete and two continuous *Actor-Critic* agents can be used, in addition to those provided through *Stable-Baselines*, see the next section.

- *Stable-Baselines*: Stable-Baselines provides a number of open-source implementations for various Reinforcement-Learning algorithms. In our framework, we use the latest version of these algorithms, through *Stable-Baselines3* [Raf+21]. The advantage when using algorithms provided through Stable-Baselines lies in the fact that they need close to no custom implementation from our end, we can instead interact with them through very simple interfaces. This cuts down on the amount of time and effort that needs to be spent developing, implementing and maintaining these powerful algorithms, and allows us to introduce a higher number of algorithms than would otherwise be possible. Currently, five different *Stable-Baselines3* algorithms are used in our simulation framework, see the *Stable-Baselines3* documentation [Bas22] and the referenced papers for more information about the different algorithms:

    - *A2C*: Advantage-Actor-Critic ([Mni+16])
    - *DDPG*: Deep deterministic policy gradient ([Lil+15])
    - *PPO*: Proximal Policy Optimization ([Sch+17])
    - *SAC*: Soft Actor-Critic ([Haa+18])
    - *TD3*: Twin-delayed deep deterministic policy gradient ([FHM18])

    We will be using some of these algorithms in our experiments and briefly compare them with our Rule-Based approaches using our various monitoring tools, which are introduced in the next chapter, **4 Approaches to monitoring agents**.

4 Approaches to monitoring agents

In this section we will take a look at the different approaches we took to monitoring agents in our framework, explaining the reasons why we chose to implement specific features and how they help us in determining an agents strengths and weaknesses.

4.1 When to monitor what

Our *workflow* (which will be explained in more detail in **5 The *recommerce* workflow**) can generally be split into two parts when it comes to monitoring and evaluating agents: *during* and *after* training. When talking about the *workflow* we refer to the process of configuring and starting a training session, where a Reinforcement-Learning agent is being trained on a specific marketplace against competitors. The *workflow* also includes the subsequent collection of data used to evaluate an agent's performance. We are also introducing the term of the *complete agent* in this section, which will be used to refer to both Reinforcement-Learning agents that have completed a training run, and Rule-Based agents which do not need training.

As mentioned above, we split our monitoring tools into the following two categories:

1. **During training**: Having data available as soon as possible without having to wait for a long training session to end is crucial to an efficient workflow. Our framework enables us to collect and visualize data while a training session is still running. This allows users to always be well-informed about the currently running experiments. In some cases, when an agent's performance is severely lacking, users may want to stop a training session before it has finished, which is enabled through these monitoring tools. We also include the *Live-monitoring* tool in this category, which runs directly after a training session has finished, see **4.2.2 Live-monitoring**.

2. **On complete agents**: After a training session has finished we have a complete and final set of data available for an agent, which enables us to perform more thorough and reliable tests. These can include simulating runs of a marketplace to gather data on the agent's performance in different scenarios and against different competitors, or running a static analysis of the agent's policy in different market states. The tools available for trained agents are in the same way also usable on Rule-Based agents.

In the following sections, we will take a look at the tools our framework provides for monitoring agents, distinguishing between the two general types of monitoring mentioned above. The goal of these sections is to give a short overview of each tool, how and why they were implemented and what value they offer to the framework as a whole. We will also discuss some features that are not yet available, explaining how they could benefit the entire workflow and enrich the overall experience.

4.2 Monitoring a training session

When talking about monitoring agents during a training session, we are always referring to Reinforcement-Learning agents, as Rule-Based agents always perform the same and cannot be trained. But even though they cannot be trained, our monitoring tools listed in the next section, **4.3 Monitoring complete agents**, can still be applied to Rule-Based agents as well, as users may want to compare different Rule-Based strategies against each other, or measure the strategy's performance on a market before training a Reinforcement-Learning agent.

Monitoring agents while they are still being trained enables us to be more closely connected to the training process. Ultimately, the goal of such monitoring tools is to be able to predict the estimated 'quality' of the final trained agent as reliably as possible while the training is still going.

4.2.1 TensorBoard

The *TensorBoard* is an external tool from the Reinforcement-Learning library *Tensor-Flow* [Mar+15]. With just a few lines of code a training session can be connected to a TensorBoard instance. We are then able to pass any number of parameters and metrics we deem interesting or useful to the TensorBoard, which then offers visualizations for each of them, updating live as the training progresses. In addition to metrics specific to our market simulation, which can be found in Table A.3, the TensorBoard also visualizes a number of specifically training-related data points, such as the number of episodes simulated per second. To access these web-based visualizations, a local server needs to be started. The diagrams created using the TensorBoard are an immensely useful tool for quickly and easily recording data and offering a first rough comparison of competitors in the market. Aside from simple visualizations, TensorBoard also offers many plugins, and even enables users to write their own [Ten20]. Plugins such as the *What-If Tool* ([Wex+20], [PAI22]), which allows users to feed trained models with hypothetical situations to study their behaviour, can have a great influence on the way users interact with the TensorBoard and their machine learning models.

4.2.2 Live-monitoring

Unlike the TensorBoard, the monitoring tools summarized under the term *Live-monitoring* were completely and from the ground up built by our team. For most of the visualizations, the *matplotlib* [Hun07] library was used. The Live-monitoring tool combines two use cases: First, it creates visualizations for all data recorded during the training, similar to those provided through the TensorBoard. This needs to be done to be able to access the visualizations even after the training has concluded, as the TensorBoard relies on abstract data files for its visualizations. By taking the data we have at the end of the training and using our own visualization tool, we create two types of diagrams: Scatterplots, which contain all samples for a certain parameter (see for example Figure 6.3 (c)) and lineplots, which show smoothed data, such as it would be available in the TensorBoard (see for example Figure 6.1). Secondly, the tool simulates a market scenario identical to the one used during training an additional time. To understand why we do this, we need some additional information: During a training session, 'intermediate' models,

as we will call them, are being saved in regular intervals. These models contain the current policy of the agent and can be used the same as any other model of complete agents, the only difference being the quality of the agent, which can change over the course of a training run, both for the better and worse. These intermediate models can then be used by a range of monitoring tools available to us. Since the models only contain the current policy of an agent but not the history of states and actions preceding the model, we need to run separate simulations on these models to be able to analyse and evaluate them. For this, we utilize our *Agent-monitoring* toolset (explained in detail in **4.3.1 Agent-monitoring**). In **6.2.1 Live- and Agent-monitoring** we will discuss the results of a training session using the Live- and Agent-monitoring tools.

4.3 Monitoring complete agents

For monitoring trained Reinforcement-Learning and Rule-Based agents, we offer three major tools: The Agent-monitoring tool allows users to simulate a large number of episodes to visualize bigger trends, the Exampleprinter simulates a single episode, offering detailed insights into market states using an overview diagram, and the Policyanalyser is a static tool which can be used to analyse a vendor's reaction to different market states and competitor actions.

4.3.1 Agent-monitoring

The *Agent-monitoring* is a highly configurable tool for monitoring and evaluating different combinations of agents and marketplaces. Figure 4.1 shows how the tool works internally.

In addition to parameters provided to the marketplace and monitored agents, the following parameters can be used to configure the *Agent-monitoring* tool itself:

Episodes: This parameter decides how many independent simulations are run in sequence. At the start of each episode, the market state will be reset and randomized. Within an episode, the vendors run through a configurable amount of time steps, during each of which they set prices (depending on the chosen economy type, this can range from only one price for new items to three prices, including a rebuy price for used items) and a set number of customers interact with them.

Plot interval: A number of diagram types enable the user to view averaged or aggregated data over a period of time. The plot interval parameter decides the size of these intervals. Smaller intervals mean more accurate but also more convoluted data points. Computational time also increases with a smaller interval size.

Marketplace: Using this parameter, the user can set the marketplace on which the monitoring session will be run. Refer to **3.1 Market scenarios** for an explanation of the different available marketplaces.

Separate markets: This parameter is a boolean flag that determines the way in which the monitoring session will handle the agents given by the user. If the flag is enabled, each agent will be initialized on a separate instance of the chosen marketplace, meaning that the agents will be monitored independent of each other. This functionality takes a lot longer than if the flag were disabled, as the whole marketplace is simulated once for each agent. While it may seem like the same

@Rainer @Alex: That diagram takes up a whole page - I would like to move it to the Appendix, is that 'allowed', since it is quite integral to the tool?
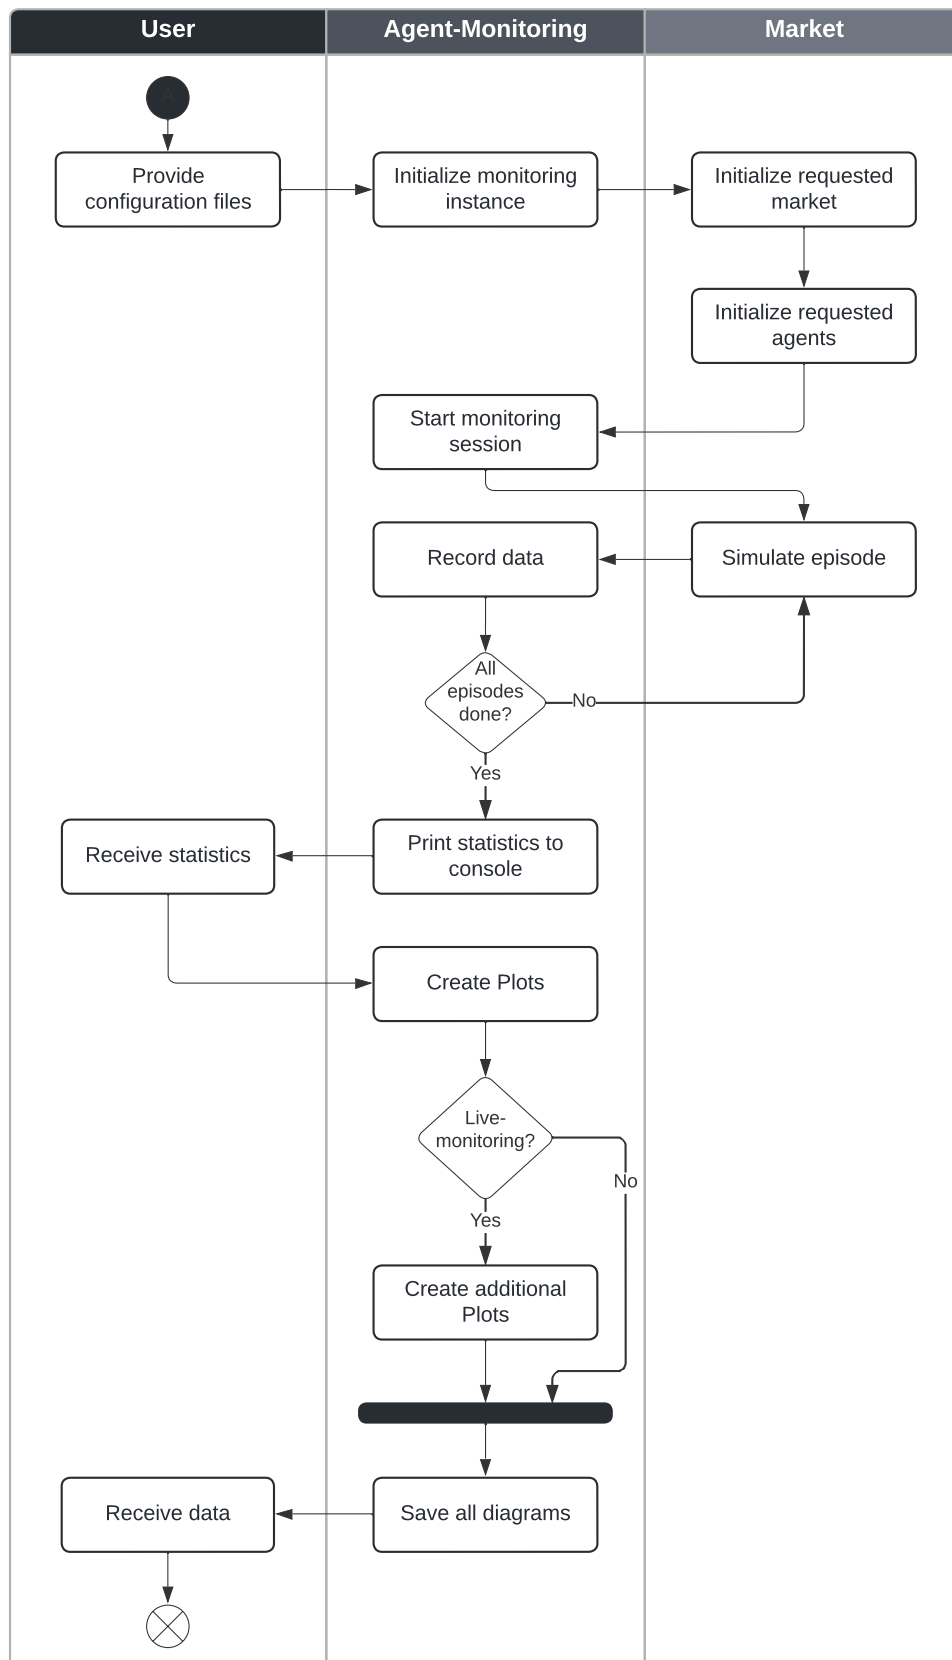
Figure 4.1: The internal workflow when running an Agent-monitoring session. Table A.3 lists the different types of diagrams created by both the Live- and Agent-Monitoring, and which of the recorded metrics they visualize.

results could be reached by simply starting multiple monitoring sessions with a different agent each, this is not the case. Using this flag instead, it is ensured that all agents get the exact same market states for each episode. By running multiple marketplaces in parallel using the *separate markets* flag, we can match the simulations as closely as possible. The most prominent use case where this flag is enabled is during the Live-monitoring after a training session, where all intermediate models are being monitored on separate markets.

If the flag is disabled, the monitoring tool will initialize only one marketplace and set the passed agents to directly compete against each other on this marketplace. This functionality is most useful when monitoring only a single agent, trying to determine its specific strengths and weaknesses against certain opponents, as it will complete a lot faster than if the flag were enabled.

Agents: Depending on the chosen marketplace, only a select number of agents can be chosen to be monitored, as each agent is built to interact with a specific type of marketplace. First off, all agents belong to one of the two major categories: *Reinforcement-Learning agent* or *Rule-Based agent* (for a more detailed overview see **3.3 Vendors**). Reinforcement-Learning agents can only be monitored on the marketplace type and market environment they were trained on, as these define the number of inputs and outputs the agent expects. Rule-Based agents can only be used on the marketplace type they were built for, as each of them makes assumptions about the number of prices they will need to set, but the market environment may be freely chosen. This leads to not all marketplace types having the same amount of Rule-Based vendors available. Following their importance for our simulation framework, the Linear Economy has the least and most often weakest vendors available, while the more refined competitors are most of the times only available as a version compatible with the Circular Economy with rebuy prices.

During each episode and for each vendor, all market events are being recorded. At the end of the monitoring session, the collected data is evaluated in different visual formats. First of all, all data that would be available to see in the *TensorBoard* during a training session is visualized using density plots. These plots can be used to compare the vendors in a more granular way, if for example the effect of a parameter on the customer's sell-back behaviour of used items should be tested. Another visualization that is created is a histogram containing the cumulative profits per episode for each agent, plotted against each other. This allows for a quick overview to see which agent had an overall better performance. One additional type of diagram is only created if the Agent-monitoring is run through the Live-monitoring tool: Violinplots. These plots, which are created for all data points available through TensorBoard (also see Table A.3), depict distributions using density curves, accentuating the minimum, maximum and median values. Violinplots are used by us to compare different training stages of an agent, as small policy changes can have great impact on these values. For exemplary Violinplots created after training, see Figure 6.4.

Aside from monitoring after a training session, a common use case of the *Agent-monitoring* tool is to test trained agents against competitors other than the ones it was trained against. This is done to test the agent's capacity to adapt to different circumstances, an important factor in deciding an agent's quality, as its competitors in the real market will differ from any it has encountered in training, due to the sheer vastness of options when it comes to dynamic pricing models available nowadays.

4.3.2 Exampleprinter

The *Exampleprinter* is a tool meant for quickly evaluating a market scenario in-depth. When run, each action taken by the monitored agents is being recorded, in addition to market states and events such as the number of customers arriving and the amount of products thrown away. At the end of this quick simulation an animated overview diagram is created, which shows all actions and their consequences for each step in the simulation, see for example Figure 6.5. Due to the large amount of data that is being collected and visualized and the overhead would come with doing so for hundreds of episodes, we chose to disconnect this functionality from large-scale tools such as the Agent-monitoring. While the Agent-monitoring tool could be seen as a tool that imitates Macro-economic behaviour, simulating hundreds of days through hundreds of episodes, the Exampleprinter instead simulates only one day, recording and visualizing all data collected during that time.

4.3.3 Policyanalyser

The last tool we want to introduce is the *Policyanalyser*. The *Policyanalyser* is our only tool which does not simulate a market. Instead, the tool can be used to monitor an agent's reaction to different market states. The user can decide on up to two different features to give as an input, such as a competitor's new and refurbished prices, and the Policyanalyser will feed all possible input combinations to the agent and record its reactions. When initializing the *Policyanalyser*, the user defines a number of parameters: The agent whose policy should be analysed, as well as the marketplace and the competitors that should be used, just as is done for all the other tools. Additionally, the user defines a **template market state**, a market state containing all values that are passed to the analysed agent, such as the number of items currently in circulation and the prices of competitors. Lastly, a list of **analysed features** needs to be provided, which defines one or two features of the template market state that should be varied. When the *Policyanalyser* is run, these features are inserted into the template market state, overwriting the initial values and creating a new combination. This new market state is then passed on to the `policy`-method of the analysed agent (for an example policy, see Figure A.2), and its reactions are recorded and visualized. See **6.2.3 Policyanalyser** for use cases of this tool.

Just mention the seciton in Appendix

    The *Policyanalyser* is the monitoring tool which operates on the smallest scale out of all the tools we built for our framework. It allows users to define any market state they want and to then accurately monitor a vendor's reactions to changes to this specific state. While the tool can just as well be utilized to test new Rule-Based strategies, it is very much meant to be used as a way to understand Reinforcement-Learning agents better, as their policies are not immediately visible to the end-user and must therefore be discovered through tools such as the one's we built.

5

The *recommerce* workflow

The main goal of the market simulation framework is to provide a simple-to-use but powerful tool for training Reinforcement-Learning algorithms on highly configurable markets for users in both a research and a business context. To achieve this, multiple components had to be developed and connected to create the workflow we now provide. This section will introduce the most important parts of the workflow.

When working with our simulation framework, users can choose from two options: First, it is possible to use our tool via a custom Command line interface (CLI). Alternatively, users can interact with the framework through a web-interface, which utilizes Docker for remote-deployment of tasks issued by the user. For detailed insights into our web-interface and remote deployment, see [Her22].

Figure 5.1 depicts the common workflow activities in our framework. For all possible tasks, the user needs to provide configuration files, which define both the task to be worked on and parameters which influence market and agent behaviour (see **5.1 Configuring the run**). After the configuration files have been validated, the simulation framework initialized the requested marketplace and agents, and then starts the requested task, for which there are currently three options provided through the CLI or web-interface: *Training*, *Agent-monitoring* and *Exampleprinter*. You may have noticed that one of our monitoring tools is missing from this list, the *Policyanalyser*. As is mentioned in **7.2.4 Policyanalyser**, this tool must unfortunately still be started manually by the user, as it has not been integrated into the rest of the workflow yet.

At the end of the respective task, the user is always provided with the various diagrams and data (e.g. trained Reinforcement-Learning models) collected during the respective task, which can then be used in subsequent tasks.

5.1 Configuring the run

Configuration is one of the most important aspects of the workflow. Without it, each simulation and training session would produce similar, if not the same results. By tweaking different parameters of a run, market dynamics can be changed and agent behaviour and thereby performance be influenced. The goal of our monitoring tools is to enable users to assess the extent to which each parameter influences certain characteristics of the training and/or monitoring session, and to enable them to make informed decisions for subsequent experiments.

Ultimately, all configuration is done using various `.json` files which contain key-value pairs of the different configurable items. We further differentiate between different groups of configurations, which means that hyperparameters influencing the market, such as maximum possible prices or storage costs, are being handled separate from parameters needed for Reinforcement-Learning agents, such as their learning rates, allowing users to easily change and tweak parameters involving different parts of the framework. Examples of such configuration files can be found in Figure A.4 and Figure A.5.
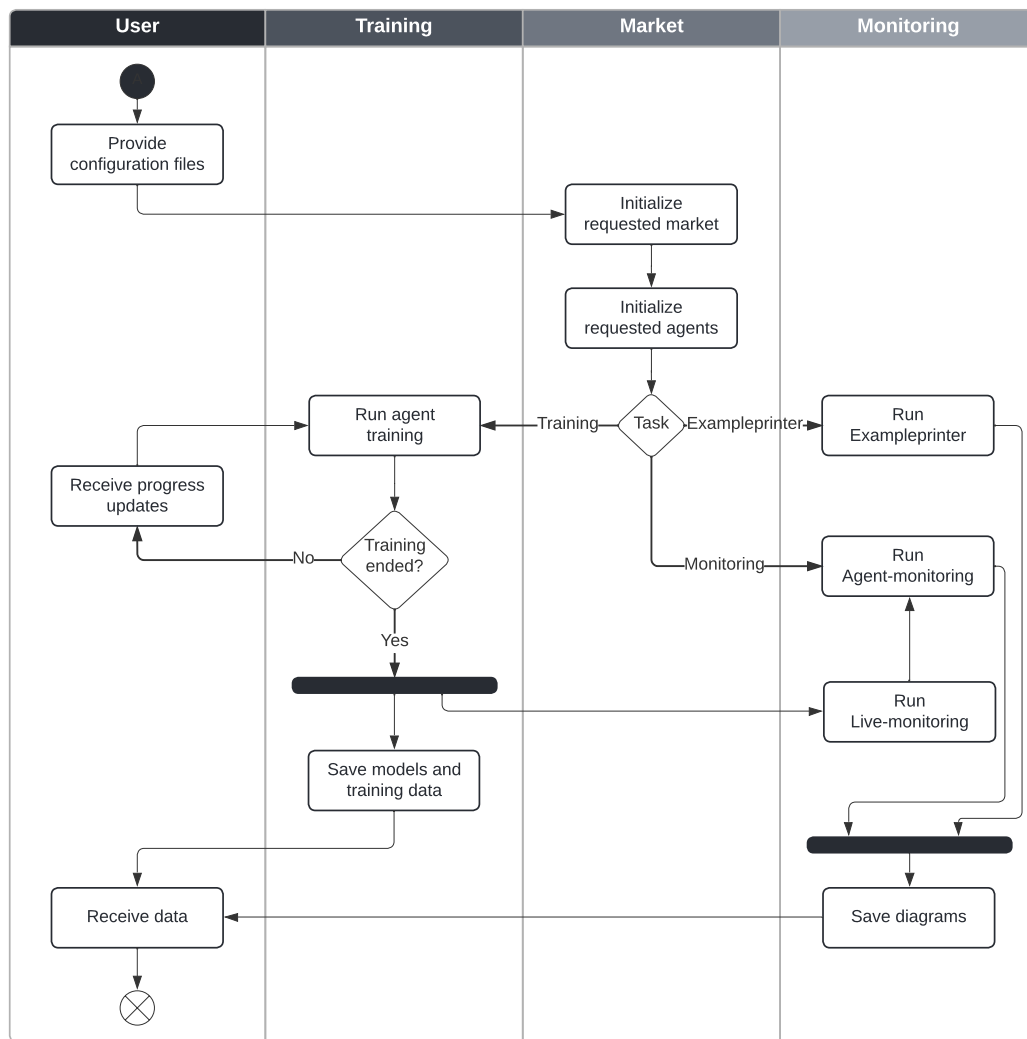
Just mention the seciton in Appendix

Figure 5.1: Diagram depicting possible workflows without webserver interaction.

5.2 The monitoring workflow

Within the recommerce workflow, there are two points in time when our monitoring tools are or can be used. While a training session is running, the TensorBoard tool is automatically used to record metrics and give insights into the current training run, by creating visualizations of current market states. By saving intermediate models, the Live-monitoring enables the user to compare agent models saved at different times within the training process, after it has concluded. This can be especially useful when trying to determine the most effective amount of training steps after which a model does not improve further, to optimize future runs. After the training has finished, or at any point disconnected from a training session, our tools described in **4.3 Monitoring complete agents** can be used to further analyse the trained models. If a training session is terminated by the user before it has completed, the Live-monitoring tool will not be run, but all intermediate models saved up to that point can still be used to monitor the agent's policy at that point in time. The next chapter, **6 Analysing Graphs**, is dedicated to the monitoring workflow, where we will first run a training session, and then monitor and evaluate the results we get from it.

6 Analysing Graphs

In this chapter, we will put the tools and workflows described in the previous sections to the use. We will train a Reinforcement-Learning agent, and then monitor it using all the tools at our disposal, highlighting where each tool is most useful, and what could be improved.

6.1 Setting up the experiment

Before starting our monitoring, we will need to conduct an experiment, where a Reinforcement-Learning algorithm is being trained on a market environment. For our experiment, we will train a Reinforcement-Learning agent using the SAC-algorithm [Haa+18] on a Duopoly marketplace with rebuy prices. The agent will be trained playing against a Rule-Based agent, more specifically the *Rule-BasedCERebuyAgentCompetitive*, as presented in **3.3.2 Data-driven models**. The configuration files for this experiment can be found in Figure A.4, Figure A.5 and Figure A.6. We will refer to this experiment as the *SAC-Duopoly* experiment. To ensure that we are evaluating an agent with a performance that is to be expected with the provided parameters, we will conduct the experiment multiple times to be able to identify outliers in the data. All diagrams except those shown in Figure 6.1 (which contains diagrams from each of the four experiment runs) have been taken from the same experiment run, denoted as SAC-Duopoly_1 in Figure 6.1 (a).

For the interested reader, a number of diagrams created through a second experiment is available in the Appendix under **A.4 PPO-Oligopoly experiment**. In this second experiment, a PPO-Agent (see **3.3.4 Machine learning models**) was trained against four Rule-Based competitors on a Circular Economy market with rebuy prices.

@Rainer @Alex: Is this expression okay, or how should I frame it?

6.2 Experiment results

In the following sections we will use our different monitoring tools on the results of the SAC-Duopoly experiment. We will start with the Live-monitoring tool, which runs automatically after the training run has completed and creates over 90 graphs and diagrams already. Due to this high number of available diagrams, we will always only look at a curated selection, highlighting those which give the best insights into the trained agent. We have also dedicated a section to diagrams which are being created, but not as useful to the user, where we will try to identify the reason for the lack of information that can be gained from these diagrams (see **6.3 Less useful diagrams**).

@Rainer @Alex: In this section I would only mention our old 'statistics plots'. For space reasons as well, should I really write that sec-

(a) SAC-Duopoly_1, 2,000 training episodes

(b) SAC-Duopoly_2, 2,000 training episodes

(c) SAC-Duopoly_3, 2,000 training episodes

(d) SAC-Duopoly_4, 6,000 training episodes

Figure 6.1: Profit per episode of four different training runs of an SAC-Agent on a Duopoly market.

6.2.1 Live- and Agent-monitoring

This section will focus on the diagrams created by the Live-monitoring tool after training, which always runs an Agent-monitoring session as well, to immediately provide the user with many additional useful diagrams without the need to run the tool manually.

A commonly asked question when deciding on the quality of a Reinforcement-Learning agent is their *stability*. If an algorithm is stable, the trained agent will produce similar results over multiple training sessions, on the condition that the parameters do not differ. Not only the rewards achieved at the end of the training will be very similar, but also the amount of episodes needed to reach certain thresholds. In the case of the SAC-Duopoly experiment, we ran the same configuration four times: The first three experiments were run using the exact same parameters, for the fourth experiment the amount of training episodes were tripled, meaning that the SAC-Agent had more time to alter its policy. Figure 6.1 shows the results of these four training sessions, created using the Live-monitoring tool and visualizes the data collected during the training process. Although many other graphs are created (Table A.3), the visualization of the total profit of the agent is the most convenient to use when evaluating an agent's performance, as the total profit is the parameter which the agent is trying to optimize.

Figure 6.1 shows the stability of the SAC-Agent very well. The agent not only reached the break-even threshold of a reward of 0 after around 150 episodes in each of the four training runs, but no matter the total length of training (see the model

in Figure 6.2 (d), which was trained three times as long as the others), the profit would always stabilize and stay at around 670. Had the monitoring tools shown that the agent performs worse than expected in some of the experiments, we could conclude that this particular algorithm is not fit for the type of work required by our market simulation.

We can also observe that the profits of the SAC-Agent and the Rule-Based agent (in the case of this experiment, a *RuleBasedCERebuyAgentStorageMinimizer*) seem to be closely linked to each other in this particular scenario. In the beginning of each experiment, when the Reinforcement-Learning agent still knows very little about the market and makes great losses, the Rule-Based agent also has a hard time to perform well. However, as soon as the SAC-Agent starts to perform better, the Rule-Based agent is also able to increase its mean profits at around the same rate as the SAC-Agent. Even more interestingly, the agents not only increase their profits at the same rate, but also very quickly arrive at a point where they earn the same mean amount as the other.

This might lead one to believe that the two competitors' policies align closely with each other. To validate this claim, another type of diagram created by the Live-monitoring tool can be used: The densityplots. These diagrams visualize probability densities for the various datasets recorded. While the diagrams shown in Figure 6.1 simply visualize data that was collected during the training run, densityplots are created by running the Agent-monitoring tool, where the marketplace is simulated an additional time. This allows us to use the 'intermediate' models we saved during the training run (see **4.2.2 Live-monitoring**) to compare the Reinforcement-Learning agent's policies at different points in time during the training.

Figure 6.2 shows the densityplots of profit-per-episode for four different training stages of the SAC-Duopoly_1 experiment. From this it can be concluded that the claim that the two competitors' policies align closely is incorrect: even though the mean reward within an episode is always very similar (Figure 6.1 (a)), the SAC-Agent achieves rewards which lie closer together, while the Rule-Based agent's rewards have more of a spread.

We can also observe that the models which were trained for longer are not necessarily better or even the same quality of the models with less experience. There is a significant improvement going from the model shown in Figure 6.2 (a) to the one in Figure 6.2 (b), this shift in the probability density curve can be correlated with the maximum mean rewards the two models could achieve during the training: For the model trained for 500 episodes, this was around 450, for the other it was about 660 (Figure 6.1 (a)). Both of these values have respectively high probabilities of being reached during the second simulation after the training has concluded, i case of the model trained for 1,000 episodes, this even coincides with the maximum in the density plot (Figure 6.2 (b)). Going from the model trained for 1,000 episodes to the next one, which was trained for 1,500 episodes (Figure 6.2 (c)), both the probability densities and the mean rewards stay very close to each other. From this we can conclude that training the SAC-Agent for more than 1,000 episodes is very likely to not have a great effect on the maximum reward achievable by the algorithm. Going from the model trained for 1,500 episodes to the last one, saved at the end of the training (Figure 6.2 (d)), we can however see a deterioration in performance: While the mean rewards hardly changed (see Figure 6.1 (a)), the probability density curve got wider at its base, extending out further towards a reward of 400, and lowering the probability of achieving a reward of 700 from previously above 0.5%
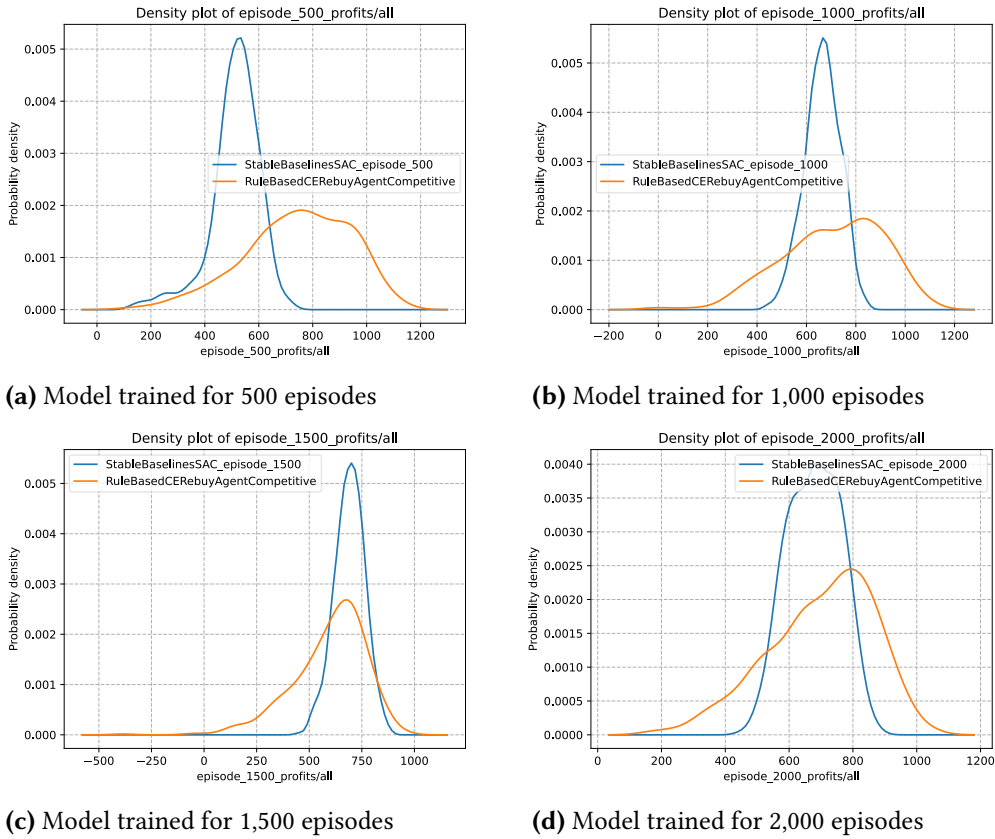
(a) Model trained for 500 episodes

(b) Model trained for 1,000 episodes

(c) Model trained for 1,500 episodes

(d) Model trained for 2,000 episodes

Figure 6.2: Probability densities for achieving a certain profit for four different training stages of the SAC-Duopoly_1 experiment.

to under 0.4%. This means that the model which was trained for the longest time produces less predictable results than those trained less, which is a tame version of so-called *Catastrophic Forgetting*, which will be explained further in **7.2.1 Live-monitoring**. The insights gained by our monitoring tools combined with the fact that we save 'intermediate' models of the Reinforcement-Learning agent during training allows us to find the optimal trained model to use for further investigation and eventual deployment in the real market. In the case of the SAC-Duopoly_1 experiment, the optimal model would be the one trained for 1,000 episodes.

6.2.1.1 Further investigation

Evaluating a model based on just the mean profits achieved by the monitored agents may not be enough for many users, which is why our Live-monitoring offers many other useful diagrams as well. In this section, we will take a look at a small selection of them. The graphs used will be from the SAC-Duopoly_1 experiment.

Users may ask themselves how the profits achieved by the different vendors are split between the two available retail channels of new and refurbished products, how many products were bought back from customers or how much the vendors had to pay for storage of these used products. For all of these questions, the Live-monitoring (together with the Agent-monitoring) tool offers two types of diagrams: First, simple lineplots as shown in Figure 6.1 can be used, as well as scatterplots which visualize the exact data recorded during the episode, instead of the trends shown by the lineplots. Figure 6.3 shows a number of different metrics, all of which
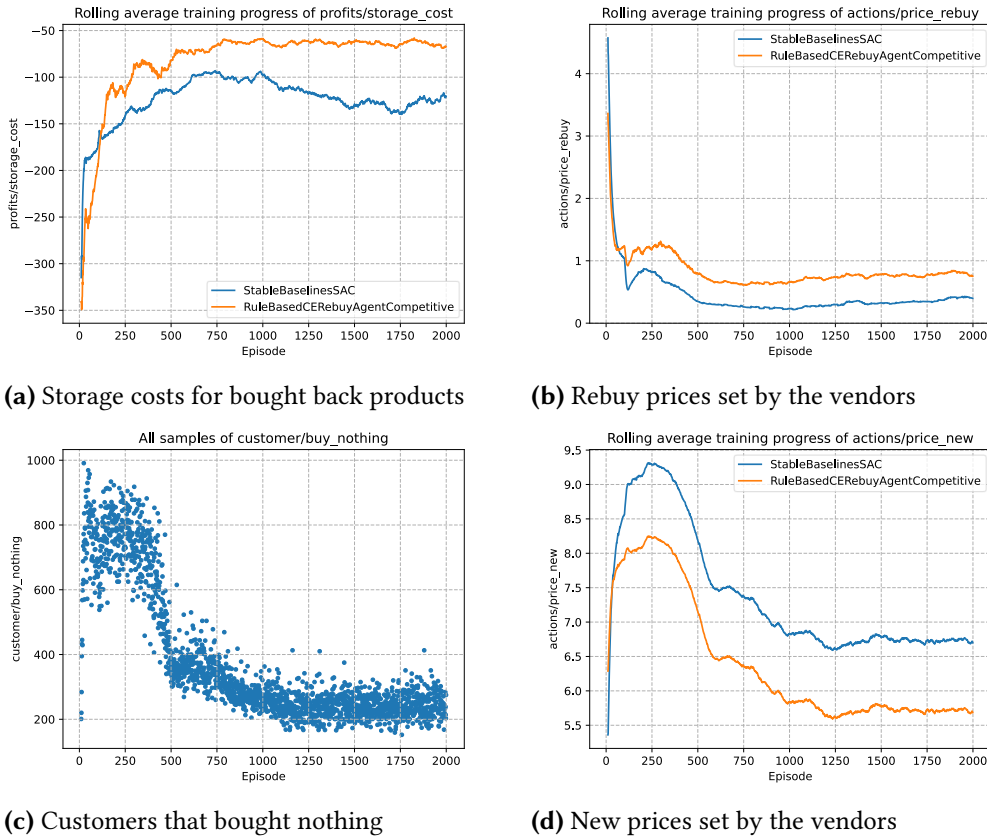
(a) Storage costs for bought back products

(b) Rebuy prices set by the vendors

(c) Customers that bought nothing

(d) New prices set by the vendors

Figure 6.3: Diagrams visualizing various data points collected during training of the SAC-Duopoly_1 experiment.

are available as both linepots and scatterplots (see Table A.3 for a list of all metrics and visualizations).

Many connections can be made when evaluating different diagrams side-by-side, such as the observation of the initially high storage costs of both vendors in Figure 6.3 (a) being caused by high rebuy prices set by the agents (Figure 6.3 (b)). High rebuy prices make it likely that customers are willing to sell back their products, which leads to (over)full inventories and high storage costs. High storage and rebuy costs will have then caused a policy change to set lower rebuy prices, thereby de-incentivizing customers to sell back as many products as they used to and lowering the agent's storage costs, increasing the profit margin.

The next pair of diagrams shows the connection between a high number of customer that buy nothing, as visualized in Figure 6.3 (c) correlating to high prices for new products, shown in Figure 6.3 (d). As mentioned in **3.2 Customers**, the customers implemented in our framework make their purchase decisions based on the prices they are presented with, which leads to many customers choosing to buy nothing before paying prices as high as set by the vendors.

If the Agent-monitoring is run after a training session (with a number of intermediate models), it also creates violinplots for all of the collected metrics (Table A.3). A selection of such diagrams can be found in Figure 6.4, showing the respective total profits (Figure 6.4 (a)) and storage costs (Figure 6.4 (b)) for each intermediate model of the trained SAC-Agent. As explained in **4.3.1 Agent-monitoring**, these plots visualize the probability distributions as shown in Figure 6.2 in a more condensed way, providing additional data such as actual maximum, minimum and median
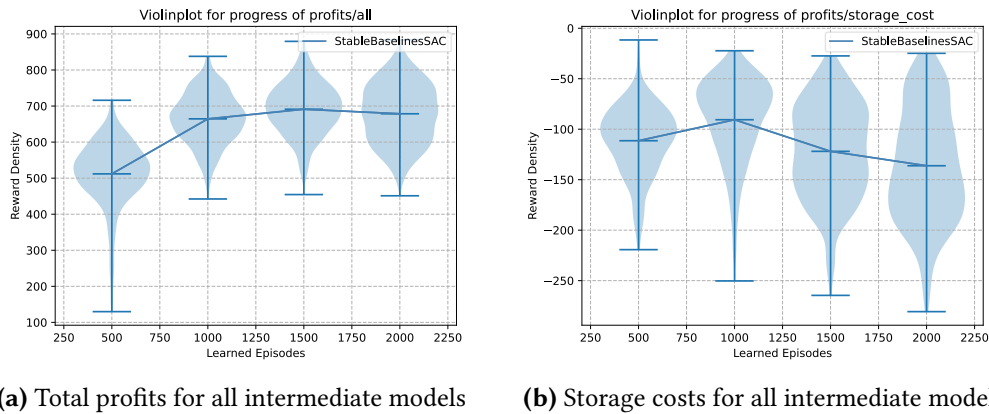
(a) Total profits for all intermediate models   **(b)** Storage costs for all intermediate models

Figure 6.4: Violinplots showing a selection of collected data when running Agent-monitoring after training.

values as well. The biggest upside of the Violinplots is however that they are able to show the distributions for all intermediate models in one diagram, which allows for even better comparisons. Looking at Figure 6.4 (a) we can immediately see the difference in the spread of total profits achieved between the models trained for 1,500 and 2,000 episodes respectively, for which we previously needed to consult and compare two different diagrams (Figure 6.2 (c) and Figure 6.2 (d)). Similarly, Figure 6.4 (b) is able to tell us something else we did not know before: the longer a model was trained for, the more likely it is that it will induce high storage costs, a trend which was not necessarily visible in Figure 6.3 (a).

Violinplots do however not replace the need for densityplots, as both have an equally useful way of displaying data. While the violinplots show rough distributions plotting against the real numbers on the y-axis, the densityplots show the concrete probability values for each possible data point.

6.2.2 Exampleprinter

All of the diagrams shown in the previous section were created as part of the automatic monitoring done after a training session has concluded. There are however two other tools at our disposal to monitor the trained Reinforcement-Learning agent after they have been trained. For this, we only need the intermediate models saved during and at the end of training. In the case of the SAC-Duopoly_1 experiment, we will be using the model saved after 1,000 episodes, as we discovered it had the best performance, see **6.2.1 Live- and Agent-monitoring**. The first tool we will now utilize is the Exampleprinter, as introduced in **4.3.2 Exampleprinter**.

To configure this monitoring run, we will again need three configuration files: one defining the task to be done (for which we will re-use the one shown in Figure A.4, only exchanging the 'task' keyword), and one for both the hyperparameter-configuration of the market and the SAC-Agent, for which we will also re-use the configuration files used for training (Figure A.5 and Figure A.6). By using the same configuration files again, we can emulate the market the agent was initially trained on as closely as possible.

At the end of the monitoring session, we receive an animated overview diagram that cycles through all time steps, allowing us to identify and examine potentially interesting time steps in the simulation. Due to the nature of this diagram being
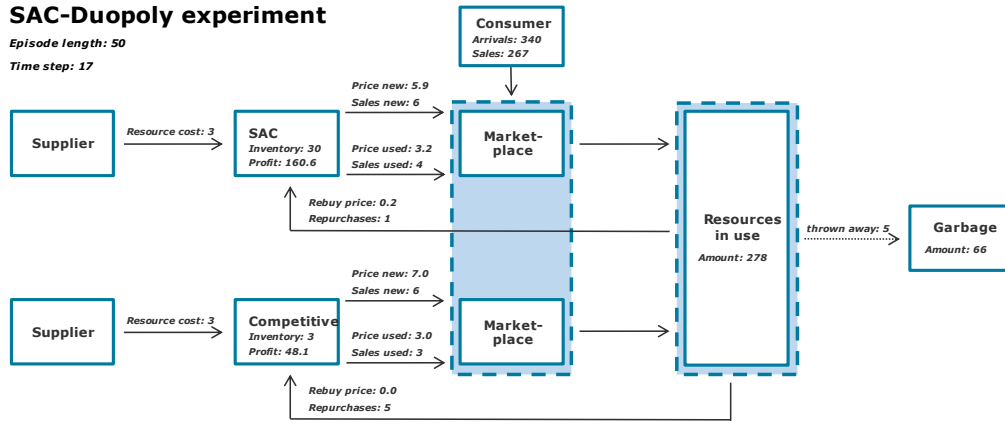
Figure 6.5: Actions and market states during step 17 of the Exampleprinter session.

animated, we will not be able to thoroughly examine the results of this monitoring tool in this thesis. Instead Figure 6.5, which shows the 17^{th} time step of the simulation, can be used to gain an understanding of the information that can be gained from this tool.

6.2.3 Policyanalyser

After a Reinforcement-Learning agent has been succesfully trained, or after implementing a new Rule-Based pricing method, users may want to analyse specific characteristics of the agent's policy. For this, the Policyanalyser can be used. We will show both use cases of the Policyanalyser, by first analysing the policy of our trained SAC-Agent and then analysing the policy of one of our Rule-Based agents, the *RuleBasedCERebuyAgentStorageMinimizer*. As was explained in **4.3.3 Policyanalyser**, we need to define both a marketplace and a template market state before running the Policyanalyser. For our experiments, we used a Circular economy with rebuy prices, and the following default market state:

- Number of products in circulation: 75

- Competitor's price for new products: 5

- Competitor's price for refurbished products: 3

- Competitor's rebuy price: 2

- Number of items in own storage: 10

- Number of items in competitor's storage: 12

In each run, the Policyanalyser will then replace two of those values, as specified by the user, with the features that should be analysed. For our first use case, we will analyse the policy of our trained SAC-Agent. The features that should be replaced are the SAC-Agent's own storage, which will range from 0 to 100, and the competitor's price for refurbished items, ranging from 0 to 10. Note that it does not matter which policy the competitor follows, as the tool does not simulate the market, but only pass market states to the monitored agent. Figure 6.6 shows the respective new and refurbished prices the SAC-Agent would set upon receiving these specific market states.
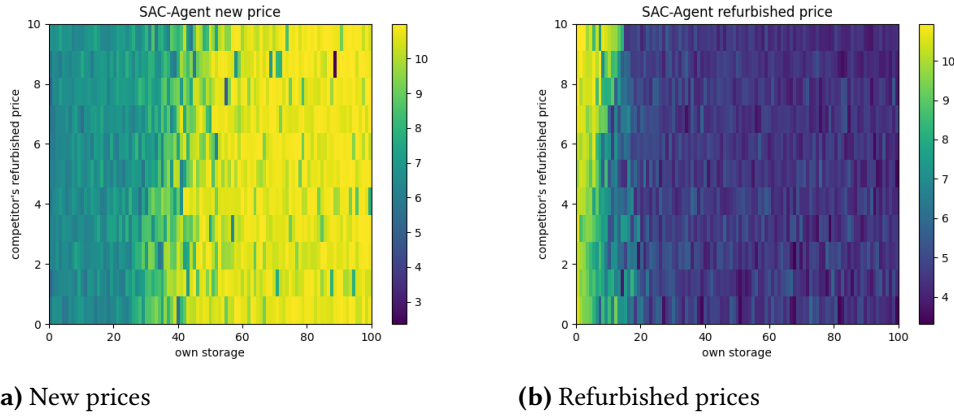
(a) New prices   **(b)** Refurbished prices

Figure 6.6: Which prices the trained SAC-Agent would set following specific market state combinations.

From Figure 6.6 (a) we can infer that the SAC-Agent is not very likely to ever set new prices lower than 6, which coincides with the observed prices set during training, see Figure 6.3 (d). Figure 6.6 (a) also shows that the SAC-Agent will increase prices for new items if it has a higher number of items in its storage, but not necessarily if the competitor's price for refurbished items rises. By increasing the price for new items, the agent tries to de-incentive customers to buy those items from it. At the same time, when looking at Figure 6.6 (b), the SAC-Agent will also decrease the price for refurbished items if it has more in storage, again to incentivize customers to buy those products. From both of these diagrams we can see that the SAC-Agent prioritizes making pricing decisions based on the amount of items in its storage, and less based on competitor prices.

Next, we will analyse the policy of the *RuleBasedCERebuyAgentStorageMinimizer* using the Policyanalyser. Figure 6.7 shows the results of this monitoring session. Figure 6.7 (a) shows the prices set by the vendor for new products, which follows the agent's policy implementation (see Figure A.2) of always undercutting competitor prices by 1, but always sell above production price (which in this case is 3, see the market configuration file in Figure A.5).

Figure 6.7 (b) shows the price the Rule-Based agent will set for buying back

(a) New prices   **(b)** Rebuy prices

Figure 6.7: Which prices the *RuleBasedCERebuyAgentStorageMinimizer* would set following specific market state combinations.

items from customers. According to its policy, the rebuy price depends on both the agent's own storage and either its own price for new items (if inventory is low) or competitor's rebuy prices, both of which can be seen nicely in the diagram. As soon as inventory exceeds the limit set in the policy, the agent sets a rebuy price that is 1 lower than the lowest competitor's rebuy price - which in this case is 2, as defined by the template market state. If inventory is low enough, rebuy prices are instead dependent on the vendor's own new prices, as seen in Figure 6.7 (a), always being one lower, also following the competitor's new price.

6.2.4 Other use cases for Agent-monitoring

Even though most of the diagrams that are created through the Agent-monitoring tool are created if it is run immediately after a training session, it is disconnected from the Live-monitoring tool. There are two major reasons for this, explained in the following two sections.

6.2.4.1 Testing Rule-Based pricing methods

First, aside from training Reinforcement-Learning agents, even though this is the main focus, our market simulation framework can just as well be used to implement and test classically Rule-Based pricing methods, as we have done ourselves for our Rule-Based vendors. In order to make this workflow possible, users need to be able to use the Agent-monitoring tool, as it is the only way of running large-scale simulations of different marketplaces aside from training a Reinforcement-Learning agent, which is not an option for this use case. Figure 6.8 shows the results of an Agent-monitoring session with a *RuleBasedCERebuyAgent* playing against a *RuleBasedCERebuyAgentStorageMinimizer*, a *Data-driven model* that tries to keep the amount of products in its storage as low as possible. Its policy can be found in Figure A.2. As was already mentioned in **3.3 Vendors**, we can see that the *RuleBasedCERebuyAgent* performed significantly worse than the other Rule-Based agent (Figure 6.8 (a)). This is mainly due to the fact that it is a purely *Inventory-based model*, while the *RuleBasedCERebuyAgentStorageMinimizer* is a *Data-driven* model. Certain characteristics of the *RuleBasedCERebuyAgentStorageMinimizer* can also be seen in Figure 6.8 (b) and Figure 6.8 (c), as it sells significantly less refurbished products, but is still able to make about the same amount of profits as the simpler *RuleBasedCERebuyAgent*. This can be attributed to its quality of keeping storage costs low by minimizing inventory size.


(a) Total profits per vendor   **(b)** Refurbished products sold   **(c)** Profits by refurbished sales

Figure 6.8: Diagrams created during an Agent-monitoring session comparing two Rule-Based pricing methods.

(a) Total profits   **(b)** Storage costs   **(c)** Sales of refurbished items

Figure 6.9: Diagrams created during an Agent-monitoring with the SAC-model saved after 1,000 training episodes playing against a *RuleBasedCERebuyAgentStorageMinimizer*.

6.2.4.2 Testing trained models against different competitors

The second reason for having the Agent-monitoring as a stand-alone tool is to be able to monitor trained Reinforcement-Learning agents not only with the competitors they were trained against, but any combination of other agents, including other trained Reinforcement-Learning agents. Figure 6.9 shows this use case on the example of our trained SAC-Agent from the SAC-Duopoly experiment playing not against the *RuleBasedCERebuyAgentCompetitive* it was trained against, but instead playing against a *RuleBasedCERebuyAgentStorageMinimizer*, as introduced in the previous section.

From Figure 6.9 (a) we can infer that the SAC-Agent we trained does not play as well against the *RuleBasedCERebuyAgentStorageMinimizer* as it did against the *RuleBasedCERebuyAgentCompetitive* (Figure 6.1 (a)), as it makes slightly less profits overall and is consistently being outperformed by the Rule-Based agent. This can be expected, as the trained agent no longer learns the behaviour of its competitor, but now acts according to the previously learned policy. This means that with a change of the competitor, the learned policy may not be as good as it was against the previous opponent.

The impact the different Rule-Based agent has on the simulation can also be seen in the difference between storage costs of both vendors between the two monitoring runs. While in Figure 6.3 (a), both vendors paid storage costs between 50 and 150, this was not the case for the second simulation, where storage costs were much lower for both vendors (Figure 6.9 (b)). As the *RuleBasedCERebuyAgentStorageMinimizer* focusses on keeping a low inventory at all times, it also buys back less items. This seems to have agreed with the SAC-Agent's policy, as the trained agent also kept fewer items in inventory in the simulation against this competitor. It was however not able to manage inventory as well as the Rule-Based agent, as both its storage costs were higher, and sales of refurbished items lower, see Figure 6.9 (c).

6.3 Less useful diagrams

See todo-note at the top.

7 Outlook

In this final chapter, we will look back at our framework and monitoring tools and give an outlook on how different parts of the framework, and the monitoring tools especially, could be improved in the future.

7.1 Modelling a realistic *recommerce* marketplace

We do not claim in any way that our simulation framework is exhaustive or complete. This chapter will focus on ways in which different parts of it could or should be improved in the future, to model a more realistic marketplace.

There are a number of parameters influencing market dynamics, small and great, which have not been modelled due to time constraints or their unpredictability in their effect on the market state or customer decisions. Seasonality of demand, customer retention, loyalty through branding and unforeseen effects on market dynamics such as global events (for example, the Covid-19 pandemic) are just some of the places where our market simulation diverges from the real market, following its inability to model these circumstances. However, all of these factors are either very rare in nature (global, unforeseen events), only applicable to a subset of markets (brand loyalty) or can initially be discarded as their effect on market dynamics is predictable, making them a lower priority than other parameters (seasonality).

The factors mentioned above all have an impact on both Linear and Circular markets, but there are also a number of parameters that are specific to Circular or even recommerce markets in particular. In the modern recommerce market, more and more consumers make their initial purchasing decisions with the product's eventual resale value already in mind [TP19]. Additionally, a great number of different motivations for choosing second-hand or refurbished products over traditional new ones can be identified. An exemplary study conducted in 2008 ([RG08]) classified such motivations into 15 different categories (see Table A.2), all of which could be used to add dimensions to customer behaviour in our simulation framework, thereby making the whole simulation more realistic.

Of course, this results in the fact that our framework can not be used out-of-the-box for any kind of market. However, great care was taken during the development process to build each part of the simulation in a modular way, so that new functionality can easily be integrated into it. Marketplaces, customers and vendors (Rule-Based as well as those using Reinforcement-Learning) have all been implemented as classes disconnected from each other, so that new variants of each can easily be added to the pool of available options to be used in experiments. The same goes for our monitoring tools, all of which have been built to work with any (valid) combination of input parameters. For reference, all of the classes shown in Figure 3.1 can be extended or replaced with ease. Please also refer to [Dre22] for more information on the modularity of our framework.

The lack of realism on certain areas of the framework does not directly impact the way that our monitoring tools work or can be used. It must however be noted

Add something about market circumstances, customers, and close with modularity argument. Section needs more structure

Number of pages: This next paragraph could be removed for space reasons

that the interpretation of results must always be based on the knowledge of these limitations.

7.2 Improving our monitoring tools

While the previous section focussed on way in which our simulation framework could be extended to be more realistic, this section will instead focus on ways in which our various monitoring tools could be improved. While all of the tools we currently have available are able to do what is expected of them, and they each have their own strengths, there are still a number of ways that the different tools could be enhanced.

7.2.1 Live-monitoring

Currently, the Live-monitoring tool works by first taking and visualizing all the data collected during a training run, and then running the Agent-monitoring tool on the saved intermediate models. The biggest downside to this is that users need to wait until the training session has finished before the Agent-monitoring is run, meaning that a lot of potential information is lost. Imagine a scenario where a training session was initialized to run for 10,000 episodes, saving intermediate models each 1,000 episodes. In the end, when the Live-monitoring tool is run, the user may find that the best model was the one saved after 3,000 episodes, meaning that a lot of time was wasted training an additional 7,000 episodes. The possibility of this happening may at first seem counterintuitive, but is a common phenomenon when training Reinforcement-Learning agents, known as *Catastrophic Forgetting* (see also [Cah11]). By improving the Live-monitoring tool with an option to **run the Agent-monitoring whenever an intermediate model is saved**, the user would be able to recognize such trends much faster and terminate the experiment at the right time. This feature could be further enhanced with a smart built-in option that **terminates the experiment for the user if a downward trend in performance is detected**. Specific thresholds for these terminations should also optionally be **set by the user**. For this, the data created during the Agent-monitoring tool would need to be saved in a machine-readable form - graphs and diagrams are not useful here.

7.2.2 Agent-monitoring

This brings us to improvements that could be made to the Agent-monitoring tool. At the moment, a lot of data is recorded when simulating the marketplace, but only graphs and diagrams are created as a result of the simulation. By **giving users the option to have data saved as (for example) `.csv` files**, users would be enabled to use the results of the simulation in other ways more easily, even for monitoring and evaluation tools completely disconnected from our own framework and the tools we provide. Reproducibility is also a major concern when it comes to evaluating simulation results, as was already mentioned in **2 Related Work** and is discussed in many papers such as [Hen+17] and [Isl+17]. In our simulation framework, with the start of a new episode the market state is always shuffled randomly, to allow Reinforcement-Learning algorithms to properly explore the

environment. This is however creating the problem of creating simulations which are currently impossible to reproduce, which could be solved by **introducing a seed-based system for shuffling market states and sharing this seed with the user**. When running a different simulation with the same seed, assuming that the marketplace type and environment stay the same, users can recreate the same random market states that are set at the start of an episode, allowing for even better and in-depth comparisons of different vendors, past a single run of the Agent-monitoring tool.

7.2.3 Exampleprinter

The Exampleprinter is a great tool for quickly monitoring and evaluating a certain market setup. However, only for one specific combination of marketplace type and market environment, a Duopoly scenario of a Circular economy with rebuy prices, the animated overview diagram is created. **Building templates and adding diagram support for more scenarios** should therefore be a priority when enhancing the Exampleprinter. Additionally, the **market-seed feature** introduced in **7.2.2 Agent-monitoring** should also be implemented for the Exampleprinter, to allow users to run a simulation multiple times to monitor possible discrepancies in agent behaviour and find outliers in the data. This could be further enhanced by adding a configuration option that allows for **more than one episode to be simulated at a time**.

As a combined enhancement for both the Agent-monitoring and the Exampleprinter, the **Exampleprinter could be integrated into the Agent-monitoring**, while still being its own tool, the same as the Agent-monitoring is integrated into the Live-monitoring.

7.2.4 Policyanalyser

The biggest and most important improvement to the Policyanalyser does not concern its concrete features, but the way users interact with it. Currently, all of the other monitoring tools are either integrated into some part of the workflow (e.g. the Live-monitoring), or can easily be started using simple commands, see **5 The *recommerce* workflow**. This is however not the case for the Policyanalyser, which must be started by going into the code itself, which is less than ideal from a user-perspective. So, **integrating the Policyanalyser into the workflow**, both by **creating a user-facing interface** for it and by **adding configuration options to start it after a training session** are features that should be a priority when continuing work on the framework. Users are also able to use the Policyanalyser to analyse a large number and combination of features, so **curating a list of useful feature combinations to be analysed** would aid many users when using this tool.

Bibliography

[Arc05]   New York Times (Archived). *As I.T. Goes, So Goes Forrester?* Accessed: 2022-
            06-21. 2005. URL: https://web.archive.org/web/20180613040854/https:
            //www.nytimes.com/2005/02/18/business/yourmoney/as-it-goes-so-goes-
            forrester.html (see page 1).

[Bas22]   Stable Baselines3. *RL Algorithms.* Accessed: 2022-06-06. 2022. URL: https:
            //stable-baselines3.readthedocs.io/en/master/guide/algos.html (see page 13).

[Bes22]   Nick Bessin. **The Marketplace of the Future: Simulation of Market
            Processes in Re-Commerce**. Hasso-Plattner-Institute, 2022 (see pages 7, 8).

[Bro+16]   Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John
            Schulman, Jie Tang and Wojciech Zaremba. **OpenAI Gym** (2016). DOI: 10.
            48550/ARXIV.1606.01540. URL: https://arxiv.org/abs/1606.01540 (see page 6).

[Cah11]   Andy Cahill. **Catastrophic forgetting in Reinforcement-Learning En-
            vironments**. MA thesis. University of Otago, 2011. URL: http://hdl.handle.
            net/10523/1765 (see page 34).

[CLH18]   Jui-Hung Chang, Yin Chung Leung and Ren-Hung Hwang. **A Survey and
            Implementation on Neural Network Visualization**. In: *2018 15th Interna-
            tional Symposium on Pervasive Systems, Algorithms and Networks (I-SPAN).*
            2018, 107–112. DOI: 10.1109/I-SPAN.2018.00026 (see page 6).

[den15]   Arnoud V. den Boer. **Dynamic pricing and learning: Historical origins,
            current research, and new directions**. *Surveys in Operations Research and
            Management Science* 20:1 (2015), 1–18. ISSN: 1876-7354. DOI: https://doi.org/
            10.1016/j.sorms.2015.03.001. URL: https://www.sciencedirect.com/science/
            article/pii/S1876735415000021 (see page 5).

[Dre22]   Leonard Dreeßen. **Pricing in the Re-Commerce Domain: Analysis of
            Pricing Strategies with an Online Market Simulation**. Hasso-Plattner-
            Institute, 2022 (see pages 7, 33).

[EIT13]   Jacqueline K Eastman, Rajesh Iyer and Stephanie P Thomas. **The impact of
            status consumption on shopping styles: An exploratory look at the
            millennial generation**. *Marketing Management Journal* 23:1 (2013), 57–73
            (see pages 8, 43).

[FHM18]   Scott Fujimoto, Herke van Hoof and David Meger. **Addressing Function
            Approximation Error in Actor-Critic Methods**. In: *Proceedings of the 35th
            International Conference on Machine Learning.* Ed. by Jennifer Dy and Andreas
            Krause. Vol. 80. Proceedings of Machine Learning Research. PMLR, Oct. 2018,
            1587–1596. URL: https://proceedings.mlr.press/v80/fujimoto18a.html (see
            page 13).

[GB22]   Torsten J. Gerpott and Jan Berends. **Competitive pricing on online mar-
            kets: a literature review**. *Journal of Revenue and Pricing Management* (June
            2022). ISSN: 1477-657X. DOI: 10.1057/s41272-022-00390-x. URL: https://doi.org/
            10.1057/s41272-022-00390-x (see page 5).

[Gee+19]   Ruben van de Geer, Arnoud V. den Boer, Christopher Bayliss, Christine S. M. Currie, Andria Ellina, Malte Esders, Alwin Haensel, Xiao Lei, Kyle D. S. Maclean, Antonio Martinez-Sykora, Asbjørn Nilsen Riseth, Fredrik Ødegaard and Simos Zachariades. **Dynamic pricing and learning with competition: insights from the dynamic pricing challenge at the 2017 INFORMS RM & pricing conference**. *Journal of Revenue and Pricing Management* 18:3 (June 2019), 185–203. ISSN: 1477-657X. DOI: 10.1057/s41272-018-00164-4. URL: https://doi.org/10.1057/s41272-018-00164-4 (see page 5).

[Gro22]   Jan Niklas Groeneveld. **A comparison of reinforcement learning algorithms for dynamic pricing in recommerce markets**. Hasso-Plattner-Institute, 2022 (see page 12).

[GWZ99]   Chris Gaskett, David Wettergreen and Alexander Zelinsky. **Q-Learning in Continuous State and Action Spaces**. In: *Advanced Topics in Artificial Intelligence*. Ed. by Norman Foo. Berlin, Heidelberg: Springer Berlin Heidelberg, 1999, 417–428. ISBN: 978-3-540-46695-6 (see page 12).

[Haa+18]   Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel and Sergey Levine. **Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor**. *CoRR* abs/1801.01290 (2018). arXiv: 1801.01290. URL: http://arxiv.org/abs/1801.01290 (see pages 13, 23).

[Hen+17]   Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup and David Meger. *Deep Reinforcement Learning that Matters*. 2017. DOI: 10.48550/ARXIV.1709.06560. URL: https://arxiv.org/abs/1709.06560 (see pages 6, 34).

[Her22]   Judith Herrmann. **Scalable Learning in the Cloud**. Hasso-Plattner-Institute, 2022 (see page 21).

[Hun07]   J. D. Hunter. **Matplotlib: A 2D graphics environment**. *Computing in Science & Engineering* 9:3 (2007), 90–95. DOI: 10.1109/MCSE.2007.55 (see page 16).

[Ins20]   Wuppertal Institut. *ReUse und Secondhand in Deutschland*. Accessed: 2022-06-21. 2020. URL: https://de.statista.com/statistik/daten/studie/1248873/umfrage/bevorzugter-kanal-fuer-den-verkauf-von-secondhand-produkten-in-deutschland/ (see page 1).

[Isl+17]   Riashat Islam, Peter Henderson, Maziar Gomrokchi and Doina Precup. **Reproducibility of Benchmarked Deep Reinforcement Learning Tasks for Continuous Control**. *CoRR* abs/1708.04133 (2017). arXiv: 1708.04133. URL: http://arxiv.org/abs/1708.04133 (see page 34).

[Jor+20]   Scott Jordan, Yash Chandak, Daniel Cohen, Mengxue Zhang and Philip Thomas. **Evaluating the Performance of Reinforcement Learning Algorithms**. In: *Proceedings of the 37th International Conference on Machine Learning*. Ed. by Hal Daumé III and Aarti Singh. Vol. 119. Proceedings of Machine Learning Research. PMLR, 13–18 Jul 2020, 4962–4973. URL: https://proceedings.mlr.press/v119/jordan20a.html (see page 6).

[KHG00]   Jeffrey O. Kephart, James E. Hanson and Amy R. Greenwald. **Dynamic pricing by software agents**. *Computer Networks* 32:6 (2000), 731–752. ISSN: 1389-1286. DOI: https://doi.org/10.1016/S1389-1286(00)00026-8. URL: https://www.sciencedirect.com/science/article/pii/S1389128600000268 (see page 5).

[KLM96]   Leslie Pack Kaelbling, Michael L. Littman and Andrew W. Moore. **Reinforcement Learning: A Survey**. *Journal of Artificial Intelligence Research* 4 (1996), 237–285. DOI: https://doi.org/10.1613/jair.301. URL: https://www.jair.org/index.php/jair/article/view/10166 (see page 12).

[KPM20]   KPMG. *Wie äußert sich bei Ihnen der Fokus auf nachhaltige Mode beim Shopping?* Accessed: 2022-06-21. 2020. URL: https://de.statista.com/statistik/daten/studie/1179997/umfrage/umfrage-unter-verbrauchern-zu-nachhaltigem-modekauf-in-deutschland/ (see page 1).

[KRH17]   Julian Kirchherr, Denise Reike and Marko Hekkert. **Conceptualizing the circular economy: An analysis of 114 definitions**. *Resources, Conservation and Recycling* 127 (2017), 221–232. ISSN: 0921-3449. DOI: https://doi.org/10.1016/j.resconrec.2017.09.005. URL: https://www.sciencedirect.com/science/article/pii/S0921344917302835 (see page 2).

[Lil+15]   Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver and Daan Wierstra. **Continuous control with deep reinforcement learning**. *arXiv preprint arXiv:1509.02971* (2015) (see page 13).

[Mar+15]   Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu and Xiaoqiang Zheng. *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems*. Software available from tensorflow.org. 2015. URL: https://www.tensorflow.org/ (see page 16).

[Mni+16]   Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Timothy P. Lillicrap, Tim Harley, David Silver and Koray Kavukcuoglu. **Asynchronous Methods for Deep Reinforcement Learning**. *CoRR* abs/1602.01783 (2016). arXiv: 1602.01783. URL: http://arxiv.org/abs/1602.01783 (see page 13).

[MPD02]   José del R. Millán, Daniele Posenato and Eric Dedieu. **Continuous-Action Q-Learning**. *Machine Learning* 49:2 (Nov. 2002), 247–265. ISSN: 1573-0565. DOI: 10.1023/A:1017788514716. URL: https://doi.org/10.1023/A:1017788514716 (see page 12).

[Mye97]   Roger B Myerson. **Game theory: analysis of conflict**. Harvard university press, 1997, 1 (see page 11).

[Nar+05]   Y. Narahari, C. V. L. Raju, K. Ravikumar and Sourabh Shah. **Dynamic pricing models for electronic business**. *Sadhana* 30:2 (Apr. 2005), 231–256. ISSN: 0973-7677. DOI: 10.1007/BF02706246. URL: https://doi.org/10.1007/BF02706246 (see page 10).

[Ove22]   Stack Overflow. *Stack Overflow Annual Developer Survey*. Accessed: 2022-06-23. 2022. URL: https://insights.stackoverflow.com/survey (see page 5).

[PAI22]   PAIR-Code. *What-If Tool*. Accessed: 2022-06-18. 2022. URL: https://pair-code.github.io/what-if-tool/ (see page 16).

[Pas+19]   Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Z. Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai and Soumith Chintala. **PyTorch: An Imperative Style, High-Performance Deep Learning Library**. *CoRR* abs/1912.01703 (2019). arXiv: 1912.01703. URL: http://arxiv.org/abs/1912.01703 (see page 12).

[Raf+21]   Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus and Noah Dormann. **Stable-Baselines3: Reliable Reinforcement Learning Implementations**. *Journal of Machine Learning Research* (2021) (see page 13).

[RG08]   Dominique Roux and Denis Guiot. **Measuring Second-Hand Shopping Motives, Antecedents and Consequences**. *Recherche et Applications en Marketing (English Edition)* 23:4 (2008), 63–91. DOI: 10.1177/205157070802300404. eprint: https://doi.org/10.1177/205157070802300404. URL: https://doi.org/10.1177/205157070802300404 (see pages 33, 44).

[SB18]   Richard S Sutton and Andrew G Barto. **Reinforcement learning: An introduction**. MIT press, 2018 (see page 13).

[Sch+17]   John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford and Oleg Klimov. **Proximal Policy Optimization Algorithms**. *CoRR* abs/1707.06347 (2017). arXiv: 1707.06347. URL: http://arxiv.org/abs/1707.06347 (see page 13).

[SK86]   George B. Sprotles and Elizabeth L. Kendall. **A Methodology for Profiling Consumers' Decision-Making Styles**. *Journal of Consumer Affairs* 20:2 (1986), 267–279. DOI: https://doi.org/10.1111/j.1745-6606.1986.tb00382.x. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1111/j.1745-6606.1986.tb00382.x. URL: https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1745-6606.1986.tb00382.x (see page 43).

[Ten20]   TensorFlow. *Developing a TensorBoard plugin*. Accessed: 2022-06-18. 2020. URL: https://github.com/tensorflow/tensorboard/blob/master/ADDING_A_PLUGIN.md (see page 16).

[TP19]   Linda Lisa Maria Turunen and Essi Pöyry. **Shopping with the resale value in mind: A study on second-hand luxury consumers**. *International Journal of Consumer Studies* 43:6 (2019), 549–556. DOI: https://doi.org/10.1111/ijcs.12539. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1111/ijcs.12539. URL: https://onlinelibrary.wiley.com/doi/abs/10.1111/ijcs.12539 (see page 33).

[Wex+20]   James Wexler, Mahima Pushkarna, Tolga Bolukbasi, Martin Wattenberg, Fernanda Viégas and Jimbo Wilson. **The What-If Tool: Interactive Probing of Machine Learning Models**. *IEEE Transactions on Visualization and Computer Graphics* 26:1 (2020), 56–65. DOI: 10.1109/TVCG.2019.2934619 (see page 16).

[Won+18]   Kanit Wongsuphasawat, Daniel Smilkov, James Wexler, Jimbo Wilson, Dandelion Mané, Doug Fritz, Dilip Krishnan, Fernanda B. Viégas and Martin Wattenberg. **Visualizing Dataflow Graphs of Deep Learning Models in TensorFlow**. *IEEE Transactions on Visualization and Computer Graphics* 24:1 (2018), 1–12. DOI: 10.1109/TVCG.2017.2744878 (see page 6).

List of Figures

A   Appendix

A.1 Tables

Consumer Characteristic	Description
Perfectionistic, High-Quality Conscious	Consumer searches carefully and systematically for the very best quality in products
Brand Conscious, 'Price = Quality'	Consumer is oriented towards buying the more expensive, well-known brands
Novelty and Fashion Conscious	Consumers who like new and innovative products and gain excitement from seeking out new things
Recreational and Shopping Conscious	Consumer finds shopping a pleasant activity and enjoys shopping just for the fun of it
Price Conscious/ Value for the Money	Consumer with a particularly high consciousness of sale prices and lower prices in general
Impulsive/ Careless	Consumer who buys on the spur of the moment and appears unconcerned about how much he/she spends
Confused by Overchoice	Consumer perceiving too many brands and stores from which to choose and experiences information overload in the market
Habitual/ Brand Loyal	Consumer who repetitively chooses the same favorite brands and stores

Table A.1: Consumer Shopping Styles, from [EIT13], including information from [SK86].

Table A.2: 15 areas of motivation toward second-hand shopping, from [RG08] (descriptions omitted).

		state/in_circulation	state/in_storage	action/price_new	action/price_refurbished	action/rebuy_price	owner/throw_away	owner_rebuys	customer/purchase_new	customer/purchase_refurbished	customer/buy_nothing	profit/rebuy_cost	profit/storage_cost	profit/by_new	profit/by_refurbished	profit/all	profit/reward
	Exampleprinter	X		X	X	X	X	X	X	X						X	
	TensorBoard	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
Live	Scatter	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	
Live	Line	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	
Agent	Density	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
Agent	Violin	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X
Agent	Line	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X	X

Table A.3: All metrics recorded during the simulation and which monitoring tools visualize them.

A.2 Rule-Based agents - Policies

```python
def policy(market_state) -> tuple:
  products_in_storage = market_state.products_in_storage
  price_refurbished = 0
  price_new = market.production_price
  rebuy_price = 0

  if products_in_storage < market.max_storage / 15:
    # try to fill up the storage immediately
    price_refurbished = market.max_price * 6/10
    price_new += market.max_price * 6/10
    rebuy_price = price_refurbished - 1

  elif products_in_storage < market.max_storage / 10:
    # try to fill up the storage
    price_refurbished = market.max_price * 5/10
    price_new += market.max_price * 5/10
    rebuy_price = price_refurbished - 2

  elif products_in_storage < market.max_storage / 8:
    # storage content is ok
    price_refurbished = market.max_price * 4/10
    price_new += market.max_price * 4/10
    rebuy_price = price_refurbished // 2

  else:
    # storage too full
    # try to get rid of some refurbished products
    price_refurbished = market.max_price * 2/10
    price_new += market.max_price * 7/10
    rebuy_price = 0

  price_new = min(9, price_new)
  return (price_refurbished, price_new, rebuy_price)
```

Figure A.1: Policy implementation of the *RuleBasedCERebuyAgent*, simplified for readability.

```python
def policy(market_state) -> tuple:
  own_storage = market_state.own_storage
  refurbished_prices, new_prices, rebuy_prices = get_competitor_prices()

  price_new = max(median(new_prices)-1, market.production_price + 1)

  if own_storage < market.max_storage / 15:
    # try to fill up the storage immediately
    price_refurbished = max(new_prices + refurbished_prices)
    rebuy_price = price_new - 1

  else:
    # storage too full, try to get rid of some refurbished products
    rebuy_price = min(rebuy_prices) - market.max_price * 10
    price_refurbished = np.quantile(refurbished_prices, 0.25)

  return clamped_prices(price_refurbished, price_new, rebuy_price)
```

Figure A.2: Policy implementation of the *RuleBasedCERebuyAgentStorageMinimizer*, simplified for readability.

```python
def policy(market_state) -> tuple:
  own_storage = market_state.own_storage
  comp_refurbished, comp_new, comp_rebuy = get_competitor_prices()
  price_new = max(min(comp_new) - 1, market.production_price + 1)

  if own_storage < market.max_storage / 15:
    # try to fill up the storage immediately
    price_refurbished = min(comp_refurbished) + 1
    rebuy_price = max(min(comp_rebuy) + 1, 2)
  elif own_storage < market.max_storage / 8:
    # storage content is ok
    rebuy_price = max(min(comp_rebuy) - 1, 0.25)
    price_refurbished = max(min(comp_refurbished) - 1, rebuy_price + 1)
  else:
    # storage too full, try to get rid of some refurbished products
    rebuy_price = max(min(comp_rebuy) - 2, 0)
    price_refurbished = max(np.quantile(comp_refurbished, 0.75) - 2, rebuy_price + 1)

  return clamped_prices(price_refurbished, price_new, rebuy_price)
```

Figure A.3: Policy implementation of the *RuleBasedCERebuyAgentCompetitive*, simplified for readability.

A.3 SAC-Duopoly experiment - Configuration files

```json
{
  "task": "training",
  "marketplace": "CircularEconomyRebuyPriceDuopoly",
  "agents": [
    {
      "name": "StableBaselinesSAC",
      "agent_class": "StableBaselinesSAC",
      "argument": ""
    },
    {
      "name": "RuleBasedCERebuyAgentCompetitive",
      "agent_class": "RuleBasedCERebuyAgentCompetitive",
      "argument": ""
    }
  ]
}
```

Figure A.4: The `environment_config.json` of the SAC-Duopoly experiment.

```json
{
  "max_storage": 100,
  "episode_length": 50,
  "max_price": 10,
  "max_quality": 50,
  "number_of_customers": 20,
  "production_price": 3,
  "storage_cost_per_product": 0.1,
  "opposite_own_state_visibility": true,
  "common_state_visibility": true,
  "reward_mixed_profit_and_difference": false
}
```

Figure A.5: The `market_config.json` of the SAC-Duopoly experiment.

```
{
  "learning_rate": 3e-4,
  "buffer_size": 1000000,
  "learning_starts": 100,
  "batch_size": 256,
  "tau": 0.005,
  "gamma": 0.99,
  "ent_coef": "auto"
}
```

Figure A.6: The configuration file for the SAC-Agent of the SAC-Duopoly experiment.

A.4 PPO-Oligopoly experiment

## A.4.1 Configuration files

A.4.1 Diagrams

Figure A.7: Mean profits achieved during the training run. The model was trained for 5,000 episodes.

Declaration of Authorship

I hereby declare that this thesis is my own unaided work. All direct or indirect sources used are acknowledged as references.

Potsdam, 24th June 2022

_____
Nikkel Mollenhauer